

# Advanced Cloud Computing

## Container Virtualization

---

Wei Wang  
CSE@HKUST  
Spring 2022



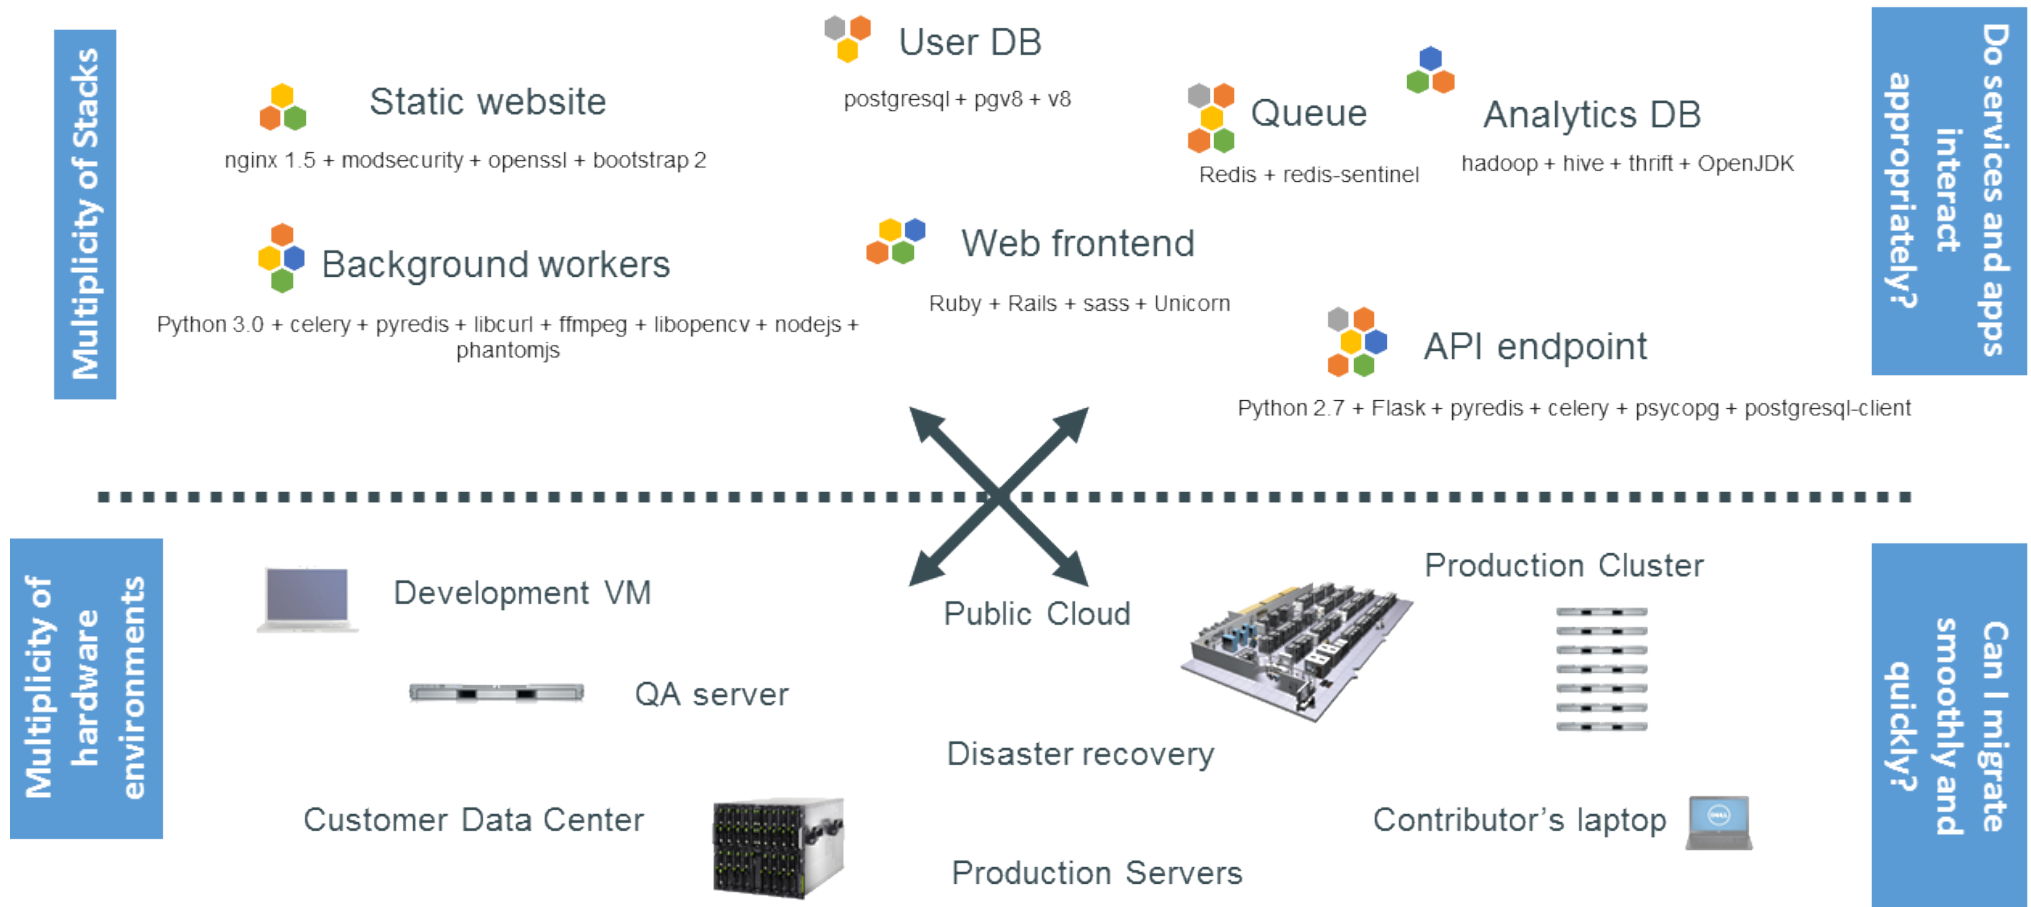
THE DEPARTMENT OF

**COMPUTER SCIENCE & ENGINEERING**








計算機科學及工程學系

# Why container?

# The challenge




# The Matrix from hell

	Static website	?	?	?	?	?	?	?
	Web frontend	?	?	?	?	?	?	?
	Background workers	?	?	?	?	?	?	?
	User DB	?	?	?	?	?	?	?
	Analytics DB	?	?	?	?	?	?	?
	Queue	?	?	?	?	?	?	?
		Development VM	QA Server	Single Prod Server	Onsite Cluster	Public Cloud	Contributor's laptop	Customer Servers
								

# Cargo transport pre-1960

Multiplicity of Goods















Do I worry about how goods interact (e.g. coffee beans next to spices)

Multiplicity of methods for transporting/storing



Can I transport quickly and smoothly (e.g. from boat to train to truck)

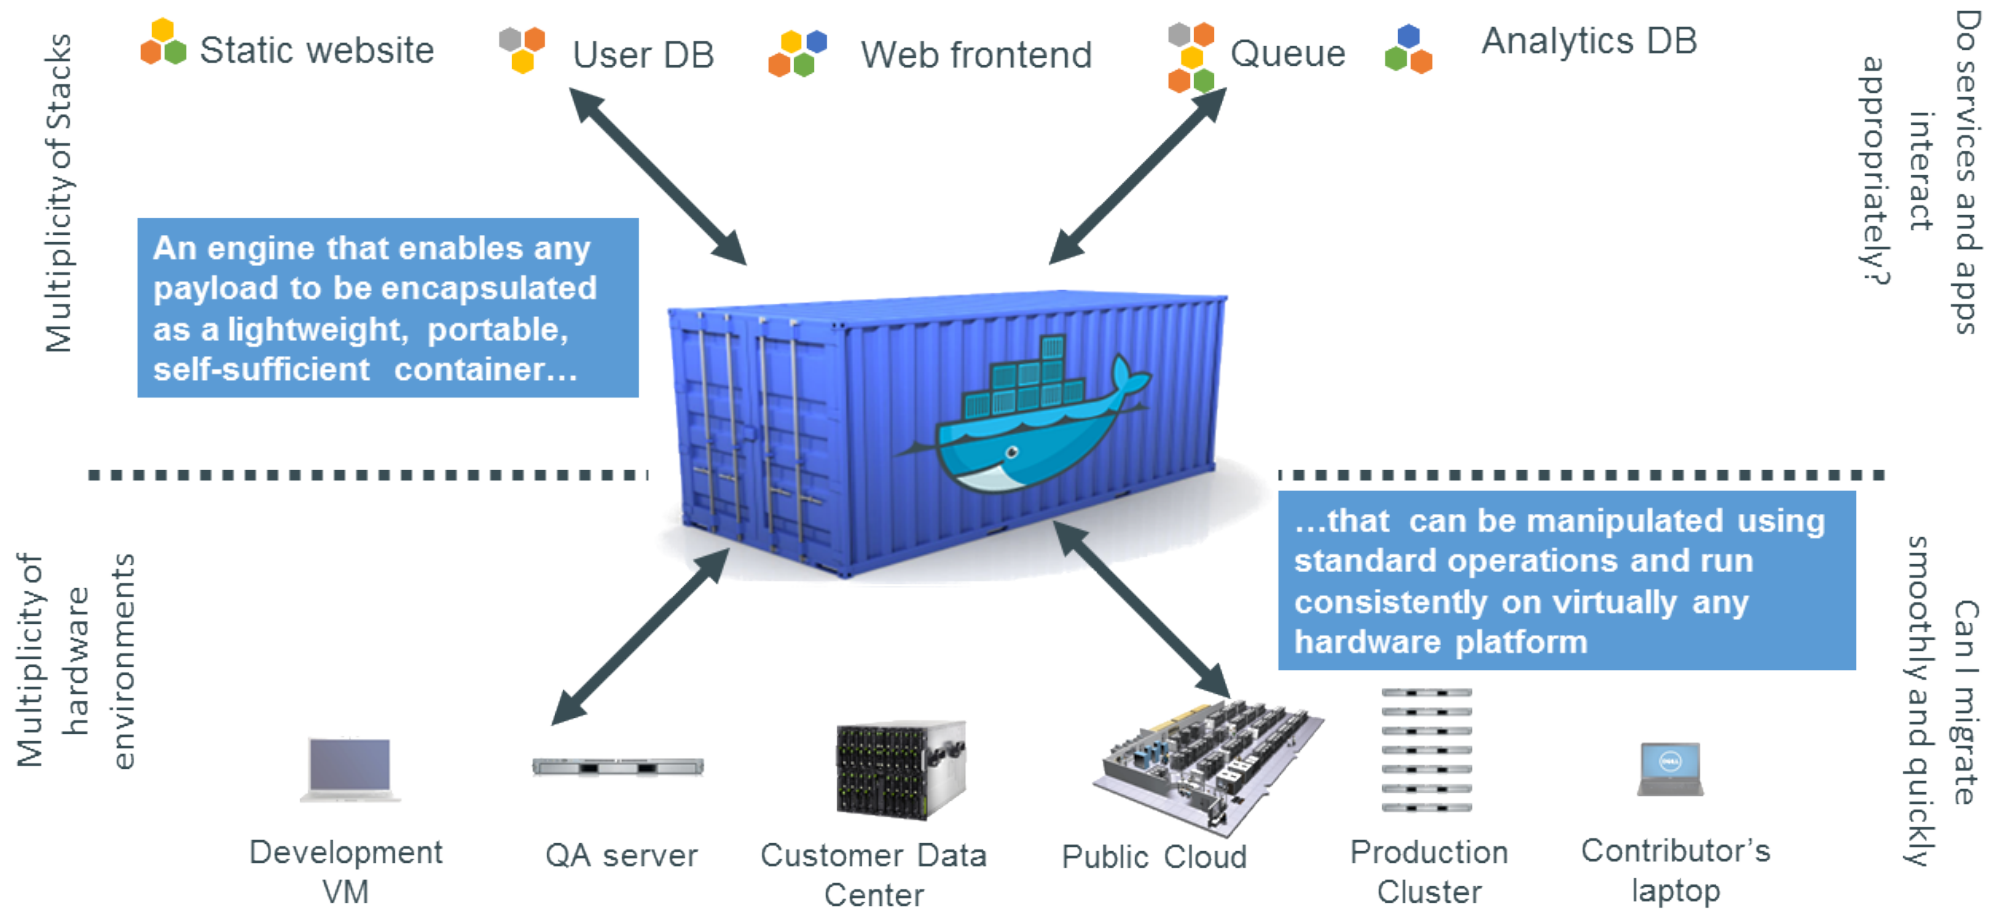
# Also a matrix from hell

	?	?	?	?	?	?	?
	?	?	?	?	?	?	?
	?	?	?	?	?	?	?
	?	?	?	?	?	?	?
	?	?	?	?	?	?	?
	?	?	?	?	?	?	?
							

# Intermodal shipping container

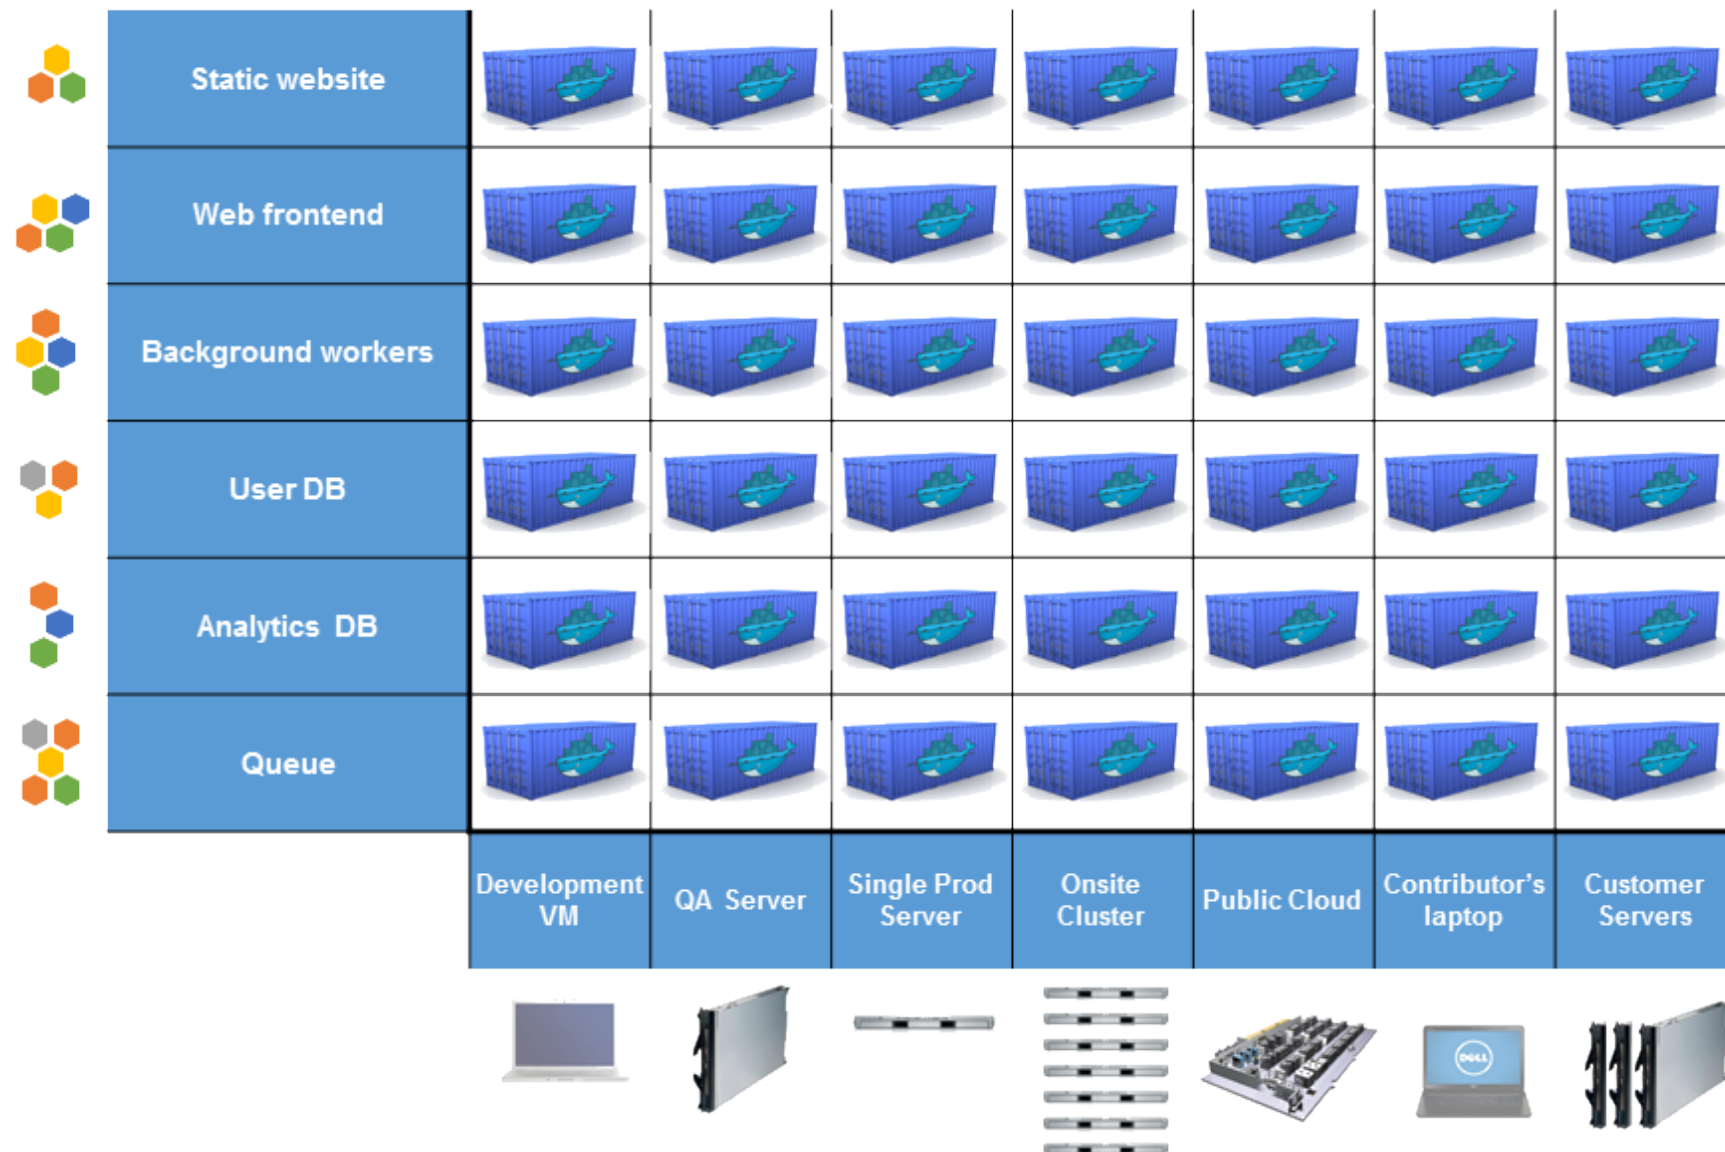


# A container system for code





# Eliminates the matrix from hell




Configure once, run anything

# Separation of concerns

- **Dan the Developer**

- Worries about what's "inside" the container
  - His code
  - His Libraries
  - His Package Manager
  - His Apps
  - His Data
- All Linux servers look the same

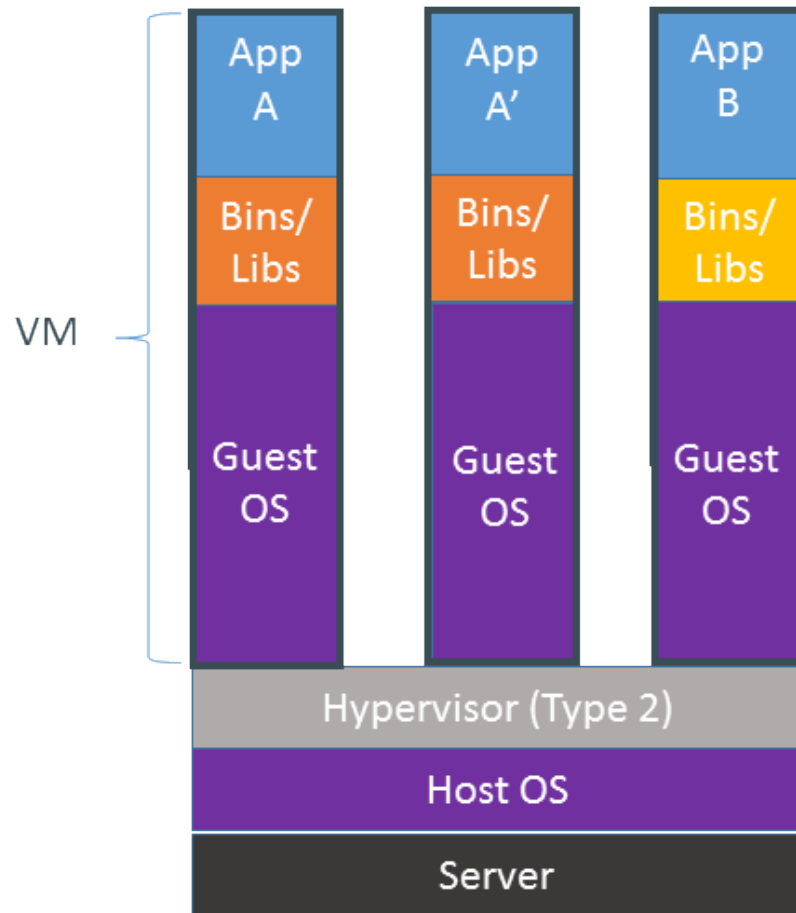


- **Oscar the Ops Guy**

- Worries about what's "outside" the container
  - Logging
  - Remote access
  - Monitoring
  - Network config
- All containers start, stop, copy, attach, migrate, etc. the same way


Configure once, run anything anywhere

# VM vs. Containers



Containers are isolated, but share OS and, where appropriate, bins/libraries

...result is significantly faster deployment, much less overhead, easier migration, faster restart



# VM vs. Container

---

## VM system call path

- ▶ application inside the VM makes a system call
- ▶ trap to the hypervisor (or host OS)
- ▶ hand trap back to the guest OS

## Container virtualization system call path

- ▶ application inside the container makes a system call
- ▶ trap to the OS
- ▶ OS returns the results to application

*No binary translation,  
no emulation*

# More technical details

---

## High level: a lightweight “VM”

- ▶ own process space
- ▶ own network interface
- ▶ can run stuffs as root
- ▶ can have its own /sbin/init (different from host)
- ▶ <<machine container>>

## Low level: chroot on steroids

- ▶ Container = isolated process: <<application container>>

# Container implementation

---

Leveraging Linux kernel mechanisms

- ▶ **namespaces**: per process resource isolation
- ▶ **cgroups**: manage resources for groups of processes
- ▶ seccomp: limit available system calls
- ▶ capabilities: limit available privileges
- ▶ CRIU: checkpoint/restore (w/ kernel support)

# What names must be virtualized?

---

## Process IDs

- ▶ **top** inside the container shows only processes running inside it
- ▶ **top** outside the container may show processes inside the container, but with different process IDs

## File names

- ▶ processes inside the container may have a limited different view of the mounted file system
- ▶ File names may resolve to different names - and some file names outside the container may be removed



# What names must be virtualized?

---

## User names

- ▶ containers may have different users w/ different roles
- ▶ **root** inside the container should not be the same as **root** outside it

## Host name and IP addresses

- ▶ processes inside the container may use a different host name and IP addresses when performing network operations

# namespaces

---

Limit the scope of kernel-side **names** and **data structures** at process granularity

<b>mnt</b>	(mount points, filesystems)	<i>CLONE_NEWNS</i>
<b>pid</b>	(processes)	<i>CLONE_NEWPID</i>
<b>net</b>	(network stack)	<i>CLONE_NEWNET</i>
<b>ipc</b>	(System V IPC)	<i>CLONE_NEWIPC</i>
<b>uts</b>	(unix timesharing - domain name, etc)	<i>CLONE_NEWUTS</i>
<b>user</b>	(UIDs)	<i>CLONE_NEWUSER</i>

Three system calls for management

- clone()** new process, new namespace, attach process to ns
- unshare()** new namespace, attach current process to it
- setns(int fd, int nstype)** join an existing namespace

# Resource control

---

The OS may want to ensure that the entire container — or everything that runs inside it — cannot consume more than a certain amount of

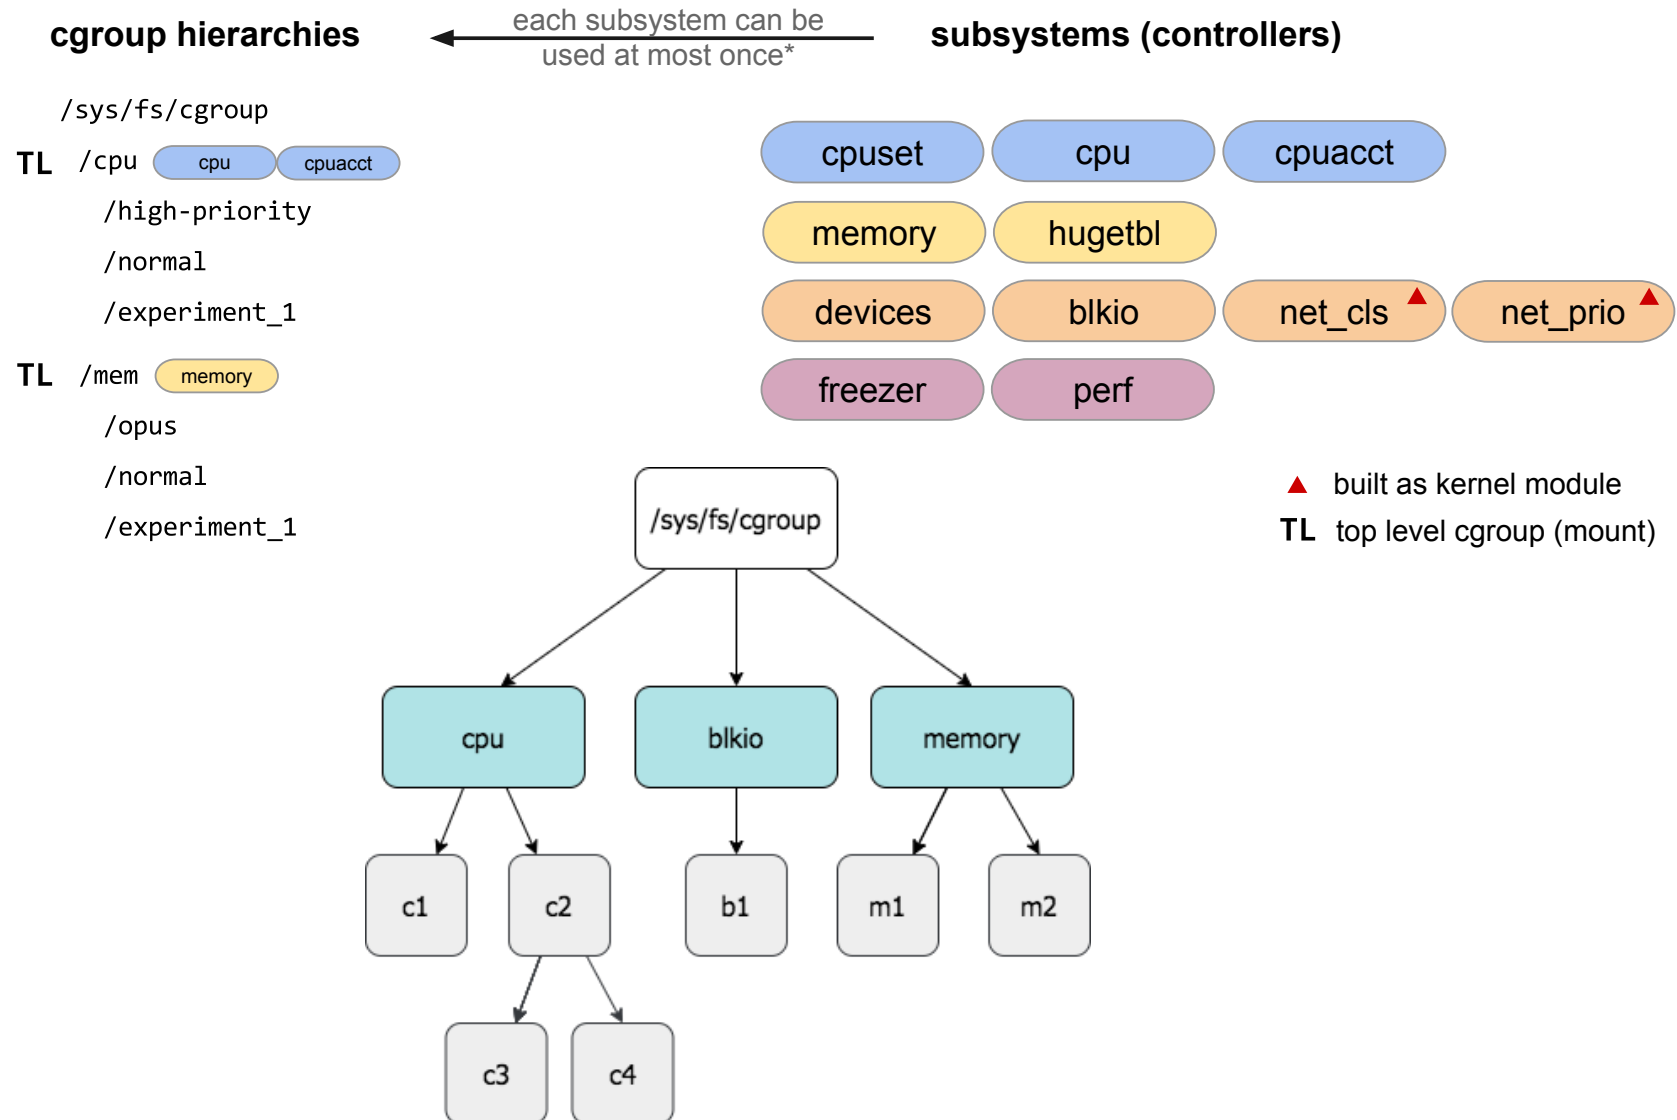
- ▶ CPU time
- ▶ memory
- ▶ disk or network bandwidth

# cgroups: Linux control groups

---

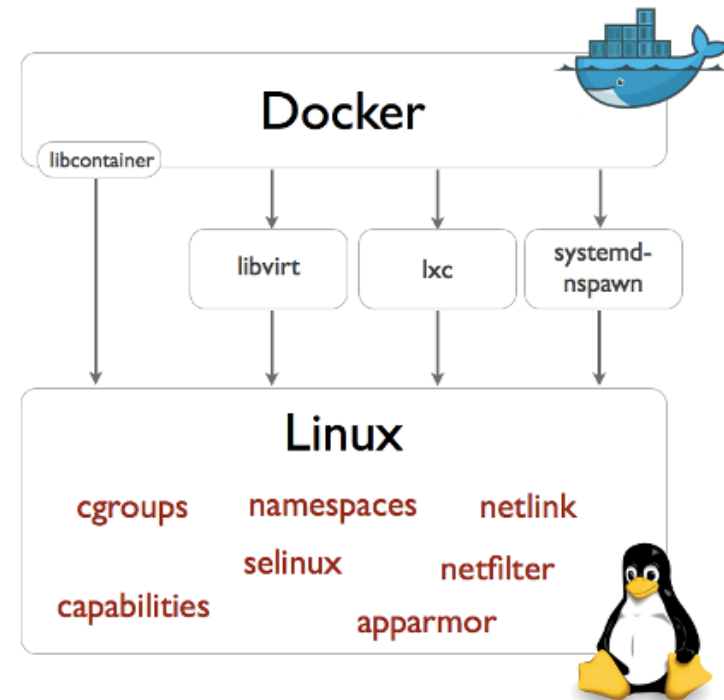
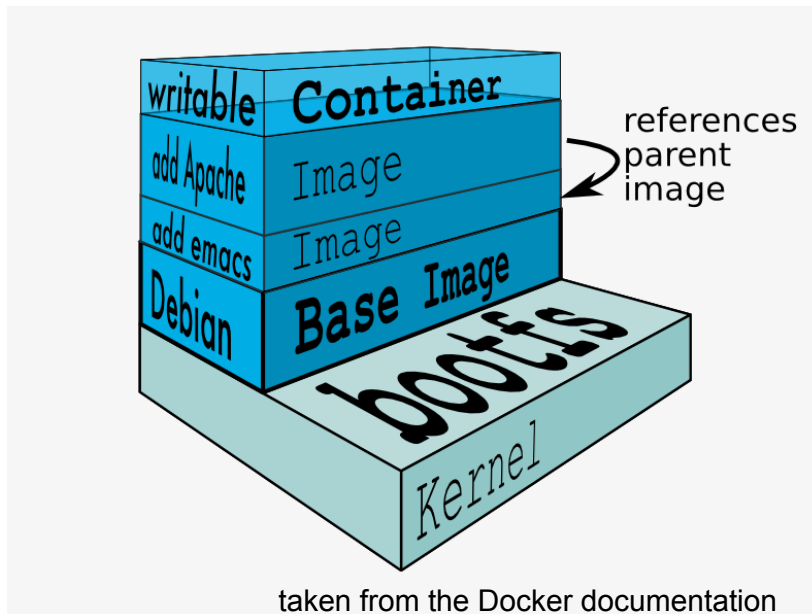
- ▶ control group subsystem offering a resource management solution for a group of processes
- ▶ Each subsystem has a *hierarchy* (a tree)
  - ▶ separate hierarchies for CPU, memory, block I/O
  - ▶ each process is in a node in each hierarchy
  - ▶ each node = a group of processes sharing the same resources

# cgroup hierarchies

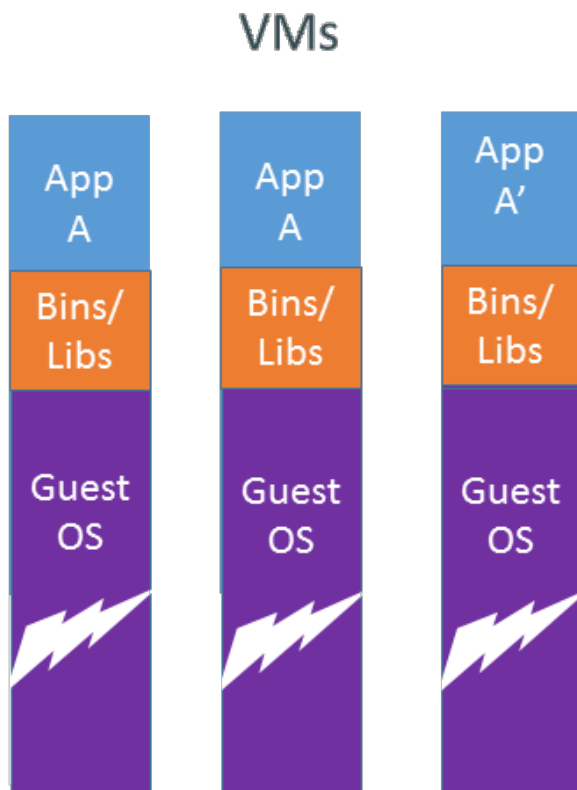


# Containers

- ▶ A light form of resource virtualization based on kernel mechanisms like **cgroups** and **namespaces**
- ▶ Multiple containers **run on the same kernel** with the illusion that they are the only one using resources

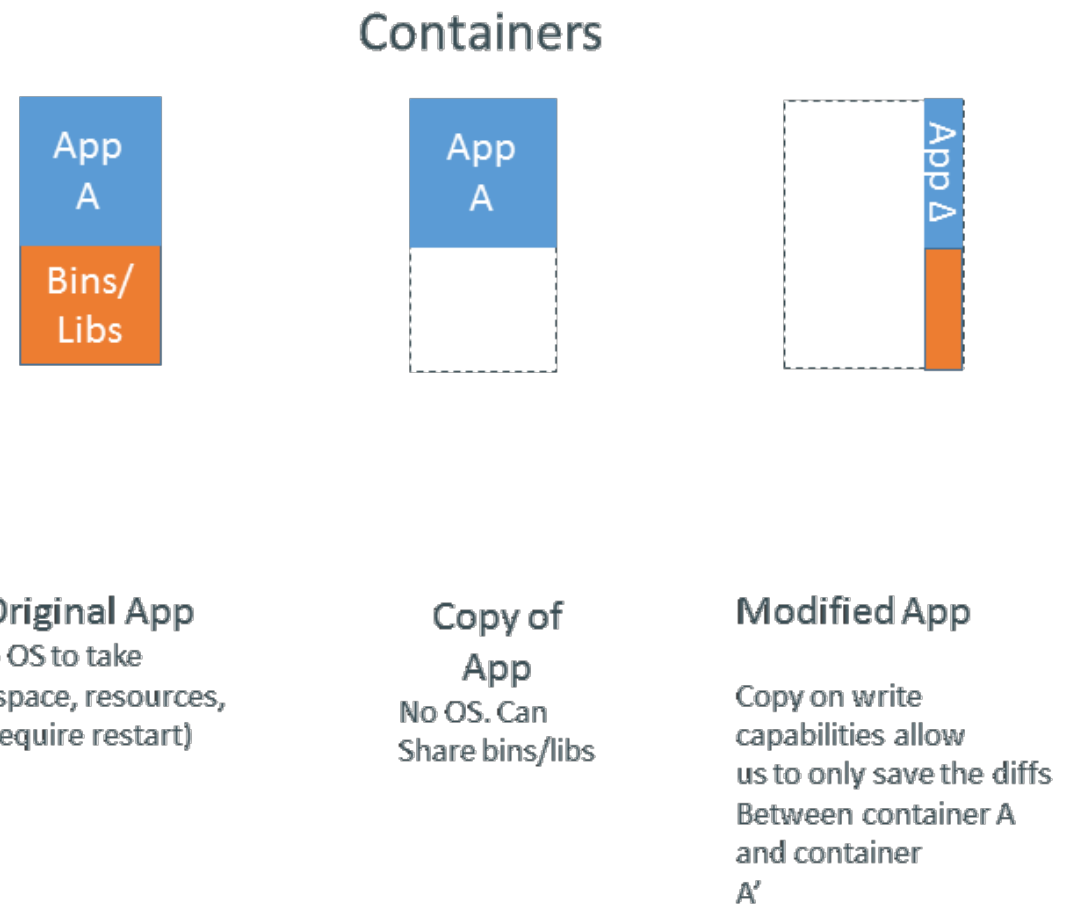


# Containers are lightweight

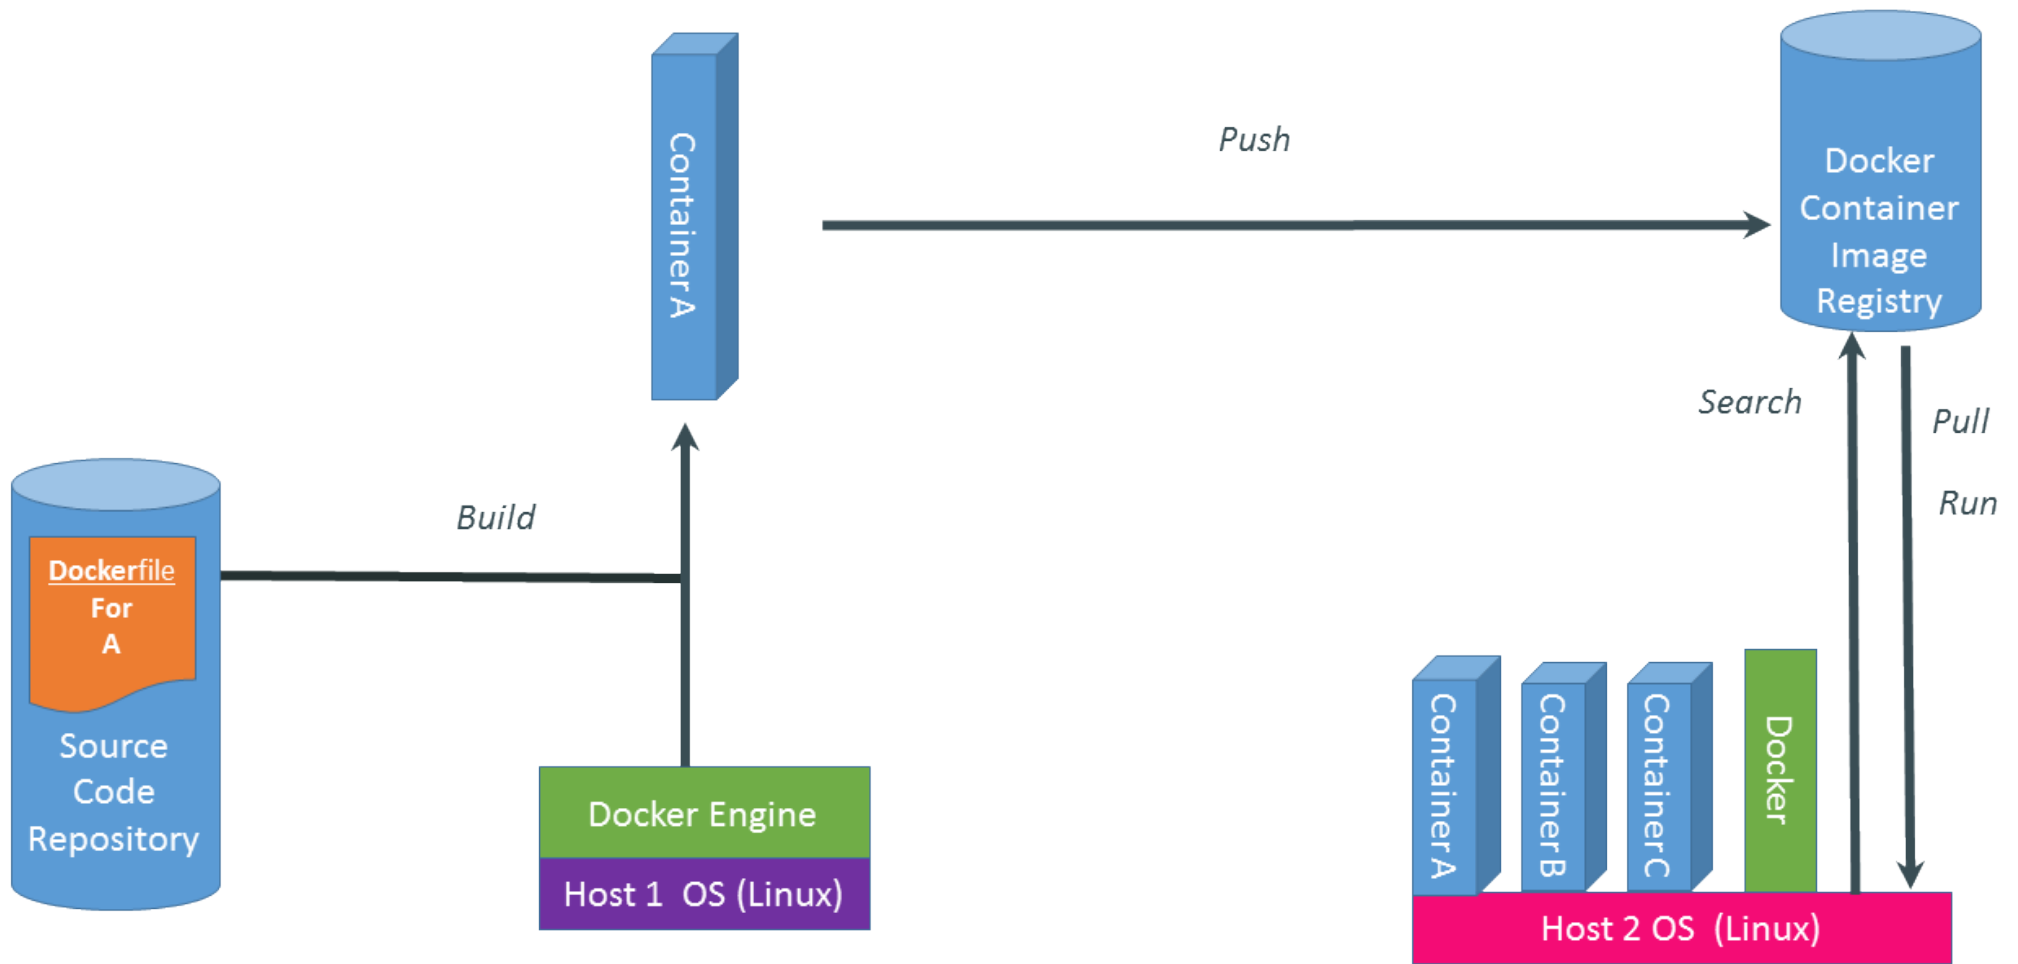


VMs

Every app, every copy of an app, and every slight modification of the app requires a new virtual server

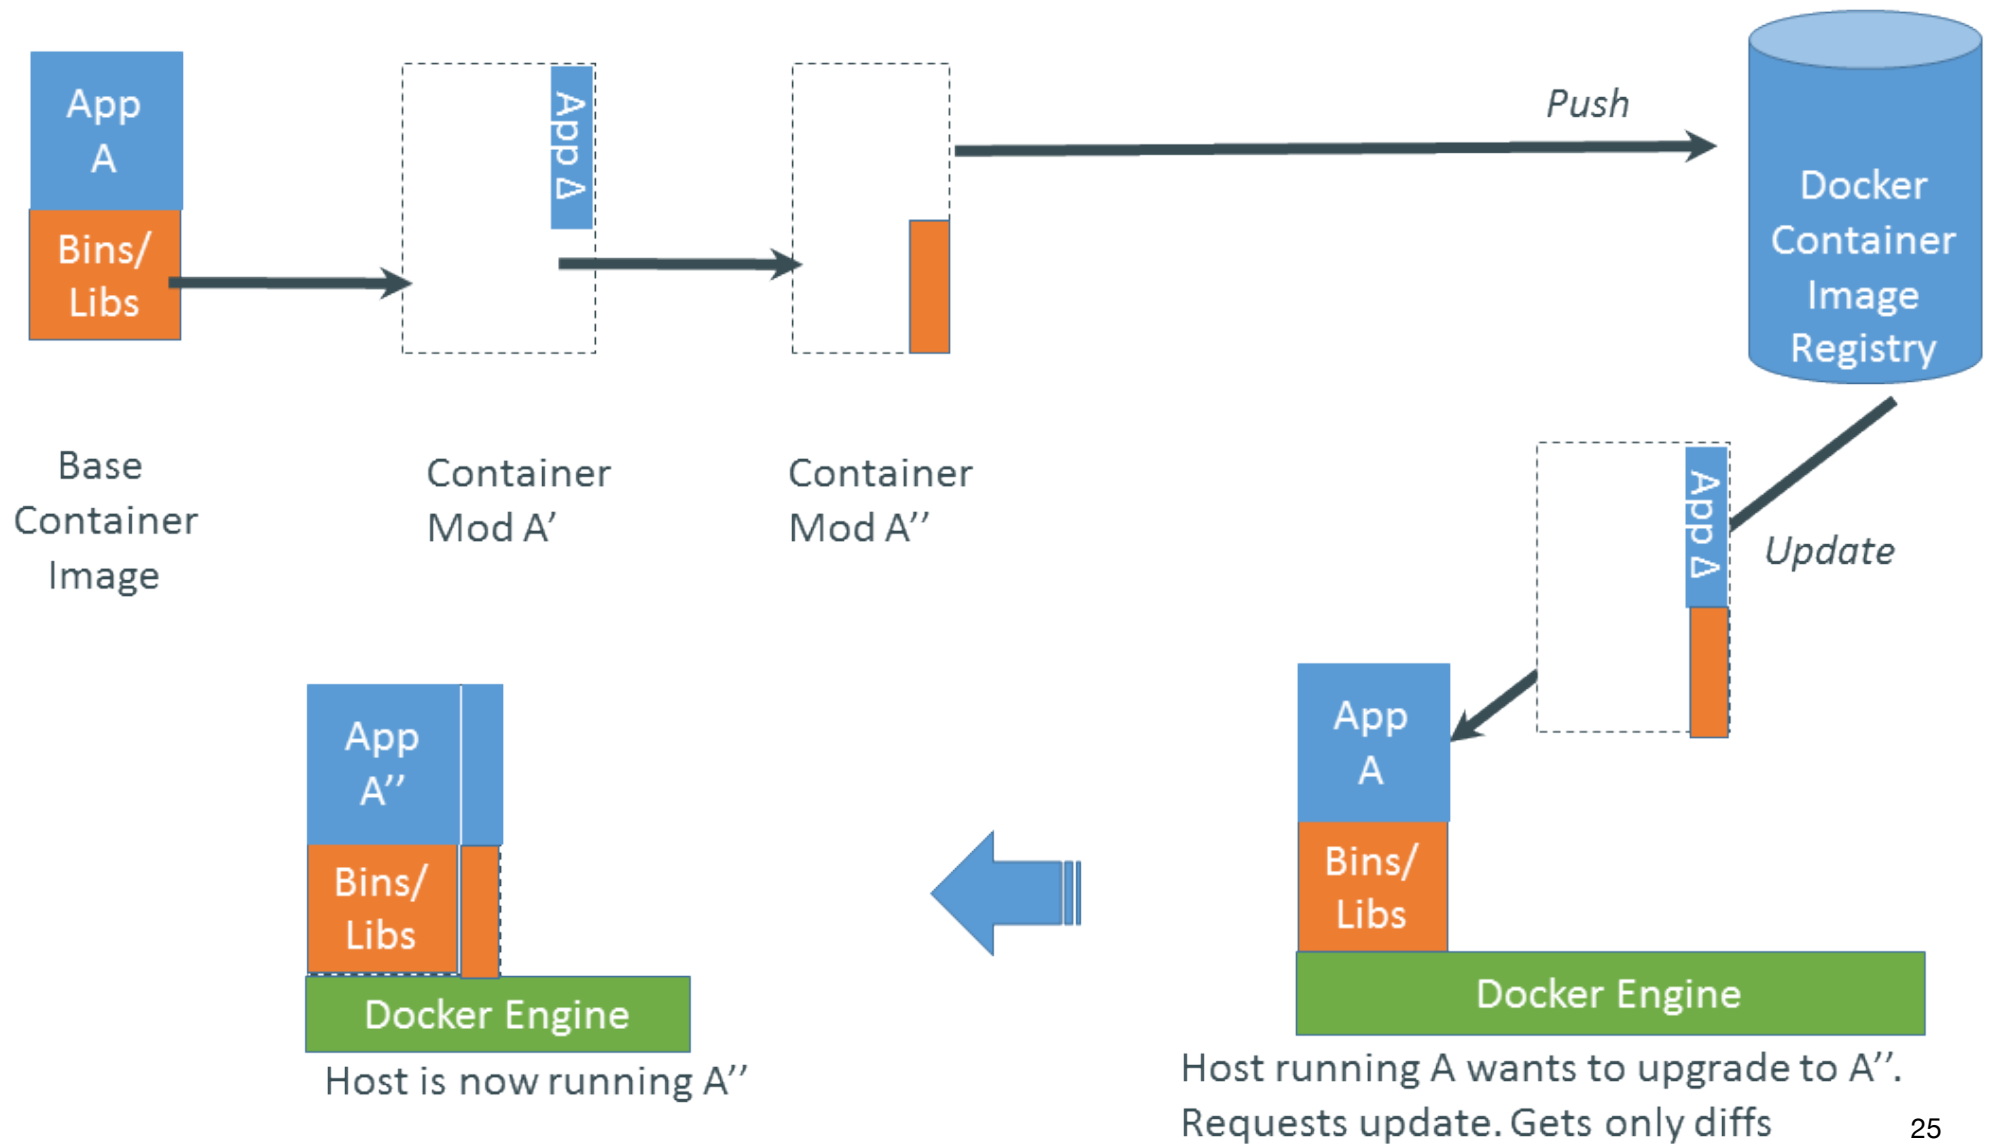


# Basics of a Docker system





# Changes and updates



# Credits

---

Slides are adapted from the community repository (CNCF) of presentations about Docker.