

Advanced Cloud Computing

Hadoop Distributed File System

Wei Wang
CSE@HKUST
Spring 2022



THE DEPARTMENT OF

COMPUTER SCIENCE & ENGINEERING

計算機科學及工程學系

Outline

HDFS overview

Architecture

Workflow

Fault tolerance

Programming APIs

What is **H**adoop **D**istributed
File **S**ystem (HDFS)?

What is HDFS?

A short answer

- ▶ An *open-source* implementation of GFS

A long answer

- ▶ A **filesystem** designed for storing **very large files** with **streaming data access patterns**, running on clusters of **commodity hardware**

Brief history

Initially developed by Doug Cutting as a filesystem for Apache Nutch, a web search engine

- ▶ early name: **N**utch **D**istributed **F**ile**S**ystem (NDFS)

Moved out of Nutch and acquired by Yahoo! in 2006 as an independent project called *Hadoop*



The origin of the name

“**Hadoop**” is a made-up name, as explained by Doug Cutting:

“The name my kid gave a stuffed yellow elephant. Short, relatively easy to spell and pronounce, meaningless, and not used elsewhere: those are my naming criteria. Kids are good at generating such. Googol is a kid’s term.”



Command-Line Interface

Basic filesystem operations

Start with `hadoop fs` (or `hdfs dfs`), followed by similar commands to Linux OS

- ▶ `% hadoop fs -ls`
- ▶ `% hadoop fs -rm`
- ▶ `% hadoop fs -mv`
- ▶ `% hadoop fs -put localFile hdfsFile`
- ▶ `% hadoop fs -get hdfsFile localFile`
- ▶ `% hadoop fs -help`

Basic filesystem operations

Also supports operations on the local filesystem (with prefix `file:///`)

- ▶ `% hadoop fs -ls file:///localDir`
- ▶ `% hadoop fs -mkdir file:///localDir`

HDFS design assumptions

Cheap, commodity machines

- ▶ Failures as a norm, rather than an exception in large clusters (e.g., 10k nodes)
- ▶ hard disk, power supply, human errors, etc.

A “modest” number of very *large* files

- ▶ a few million files each > 100MB

HDFS design assumptions

Batch processing

- ▶ Files are **write-once, mostly appended to** (perhaps concurrently)
- ▶ **Streaming reads**, rather than random data access
- ▶ High **sustained throughput** favored over low latency

Do these look familiar?

The design of HDFS heavily
borrows from GFS

Design of HDFS

NameNode

- ▶ a single master for managing filesystem meta

DataNode (chunkserver)

- ▶ multiple DataNodes for storing and retrieving data
- ▶ reports to NameNode with list of blocks hosted

SecondaryNameNode (shadow master)

- ▶ performs checkpointing

Design of HDFS

Single namespace for the entire cluster

Data coherency: Write-once, read-many-times

Files are broken up into blocks: 128MB blocks each replicated on multiple DataNodes

Intelligent client

- ▶ Client can find locations of blocks
- ▶ Client accesses data directly from DataNodes

Demo: HDFS web interface

Noticeable differences from GFS

Earlier versions have only a single writer per file

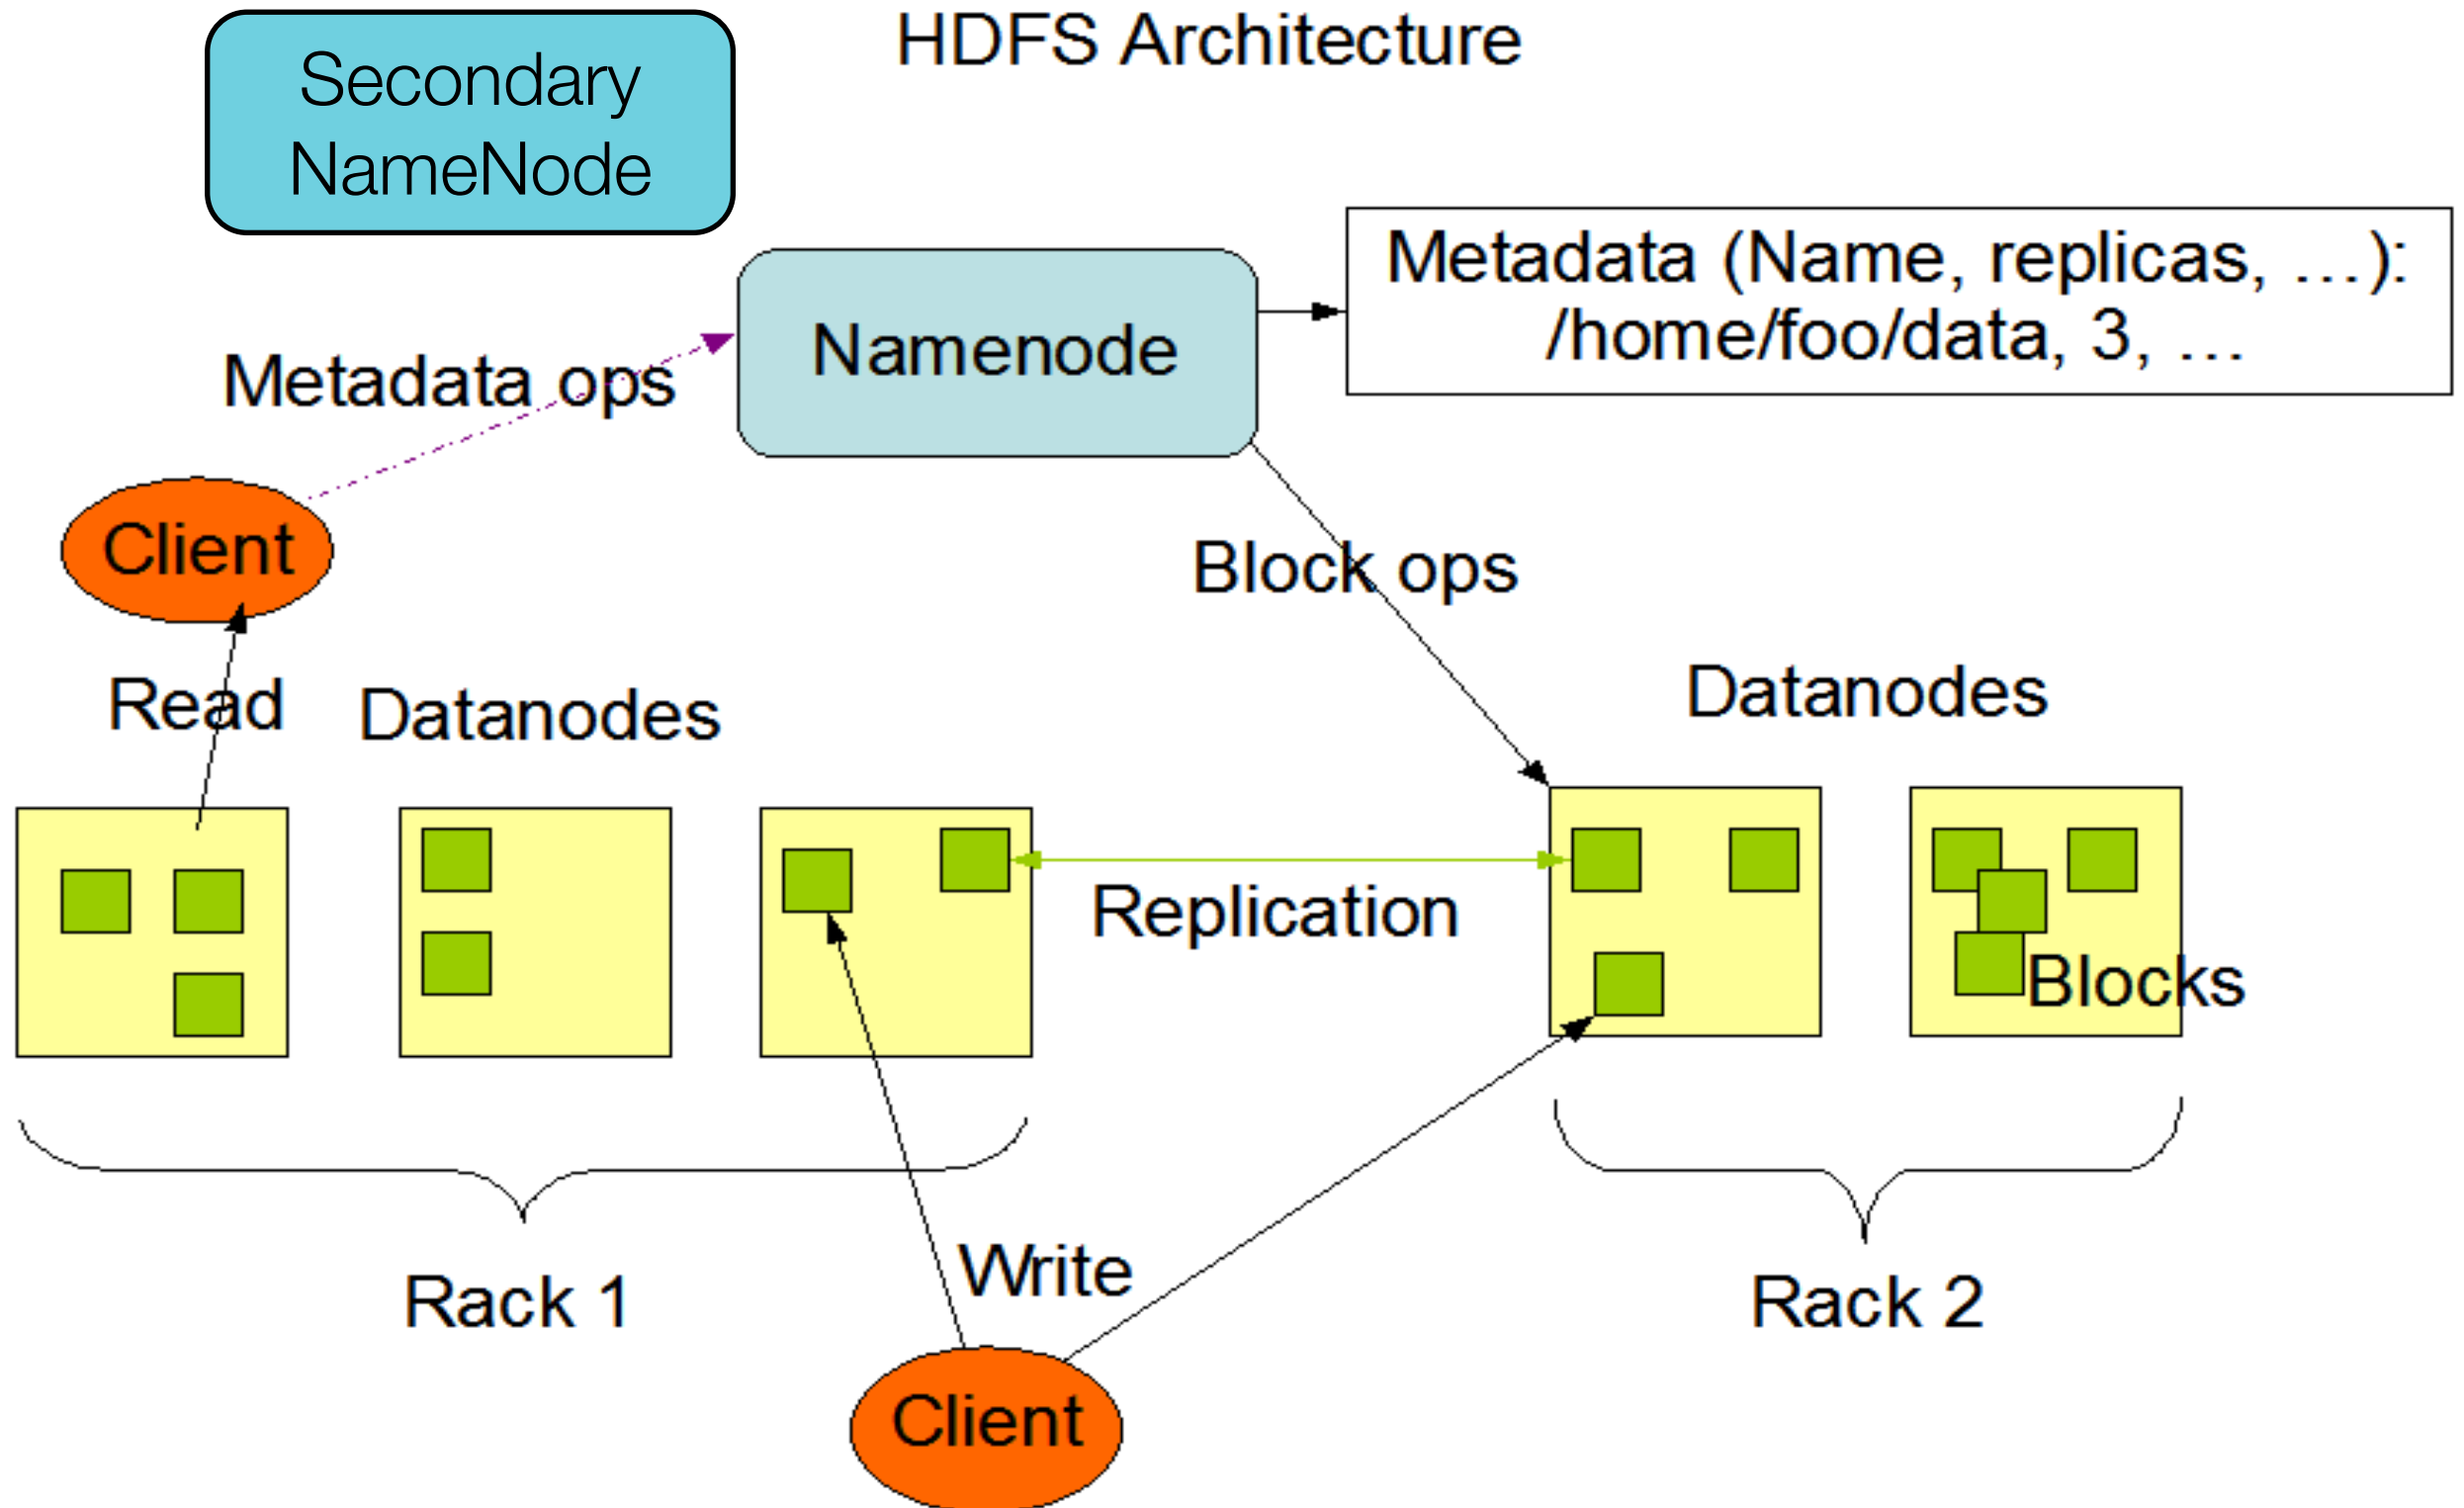
- ▶ No *record append* operation supported in earlier versions

Open source

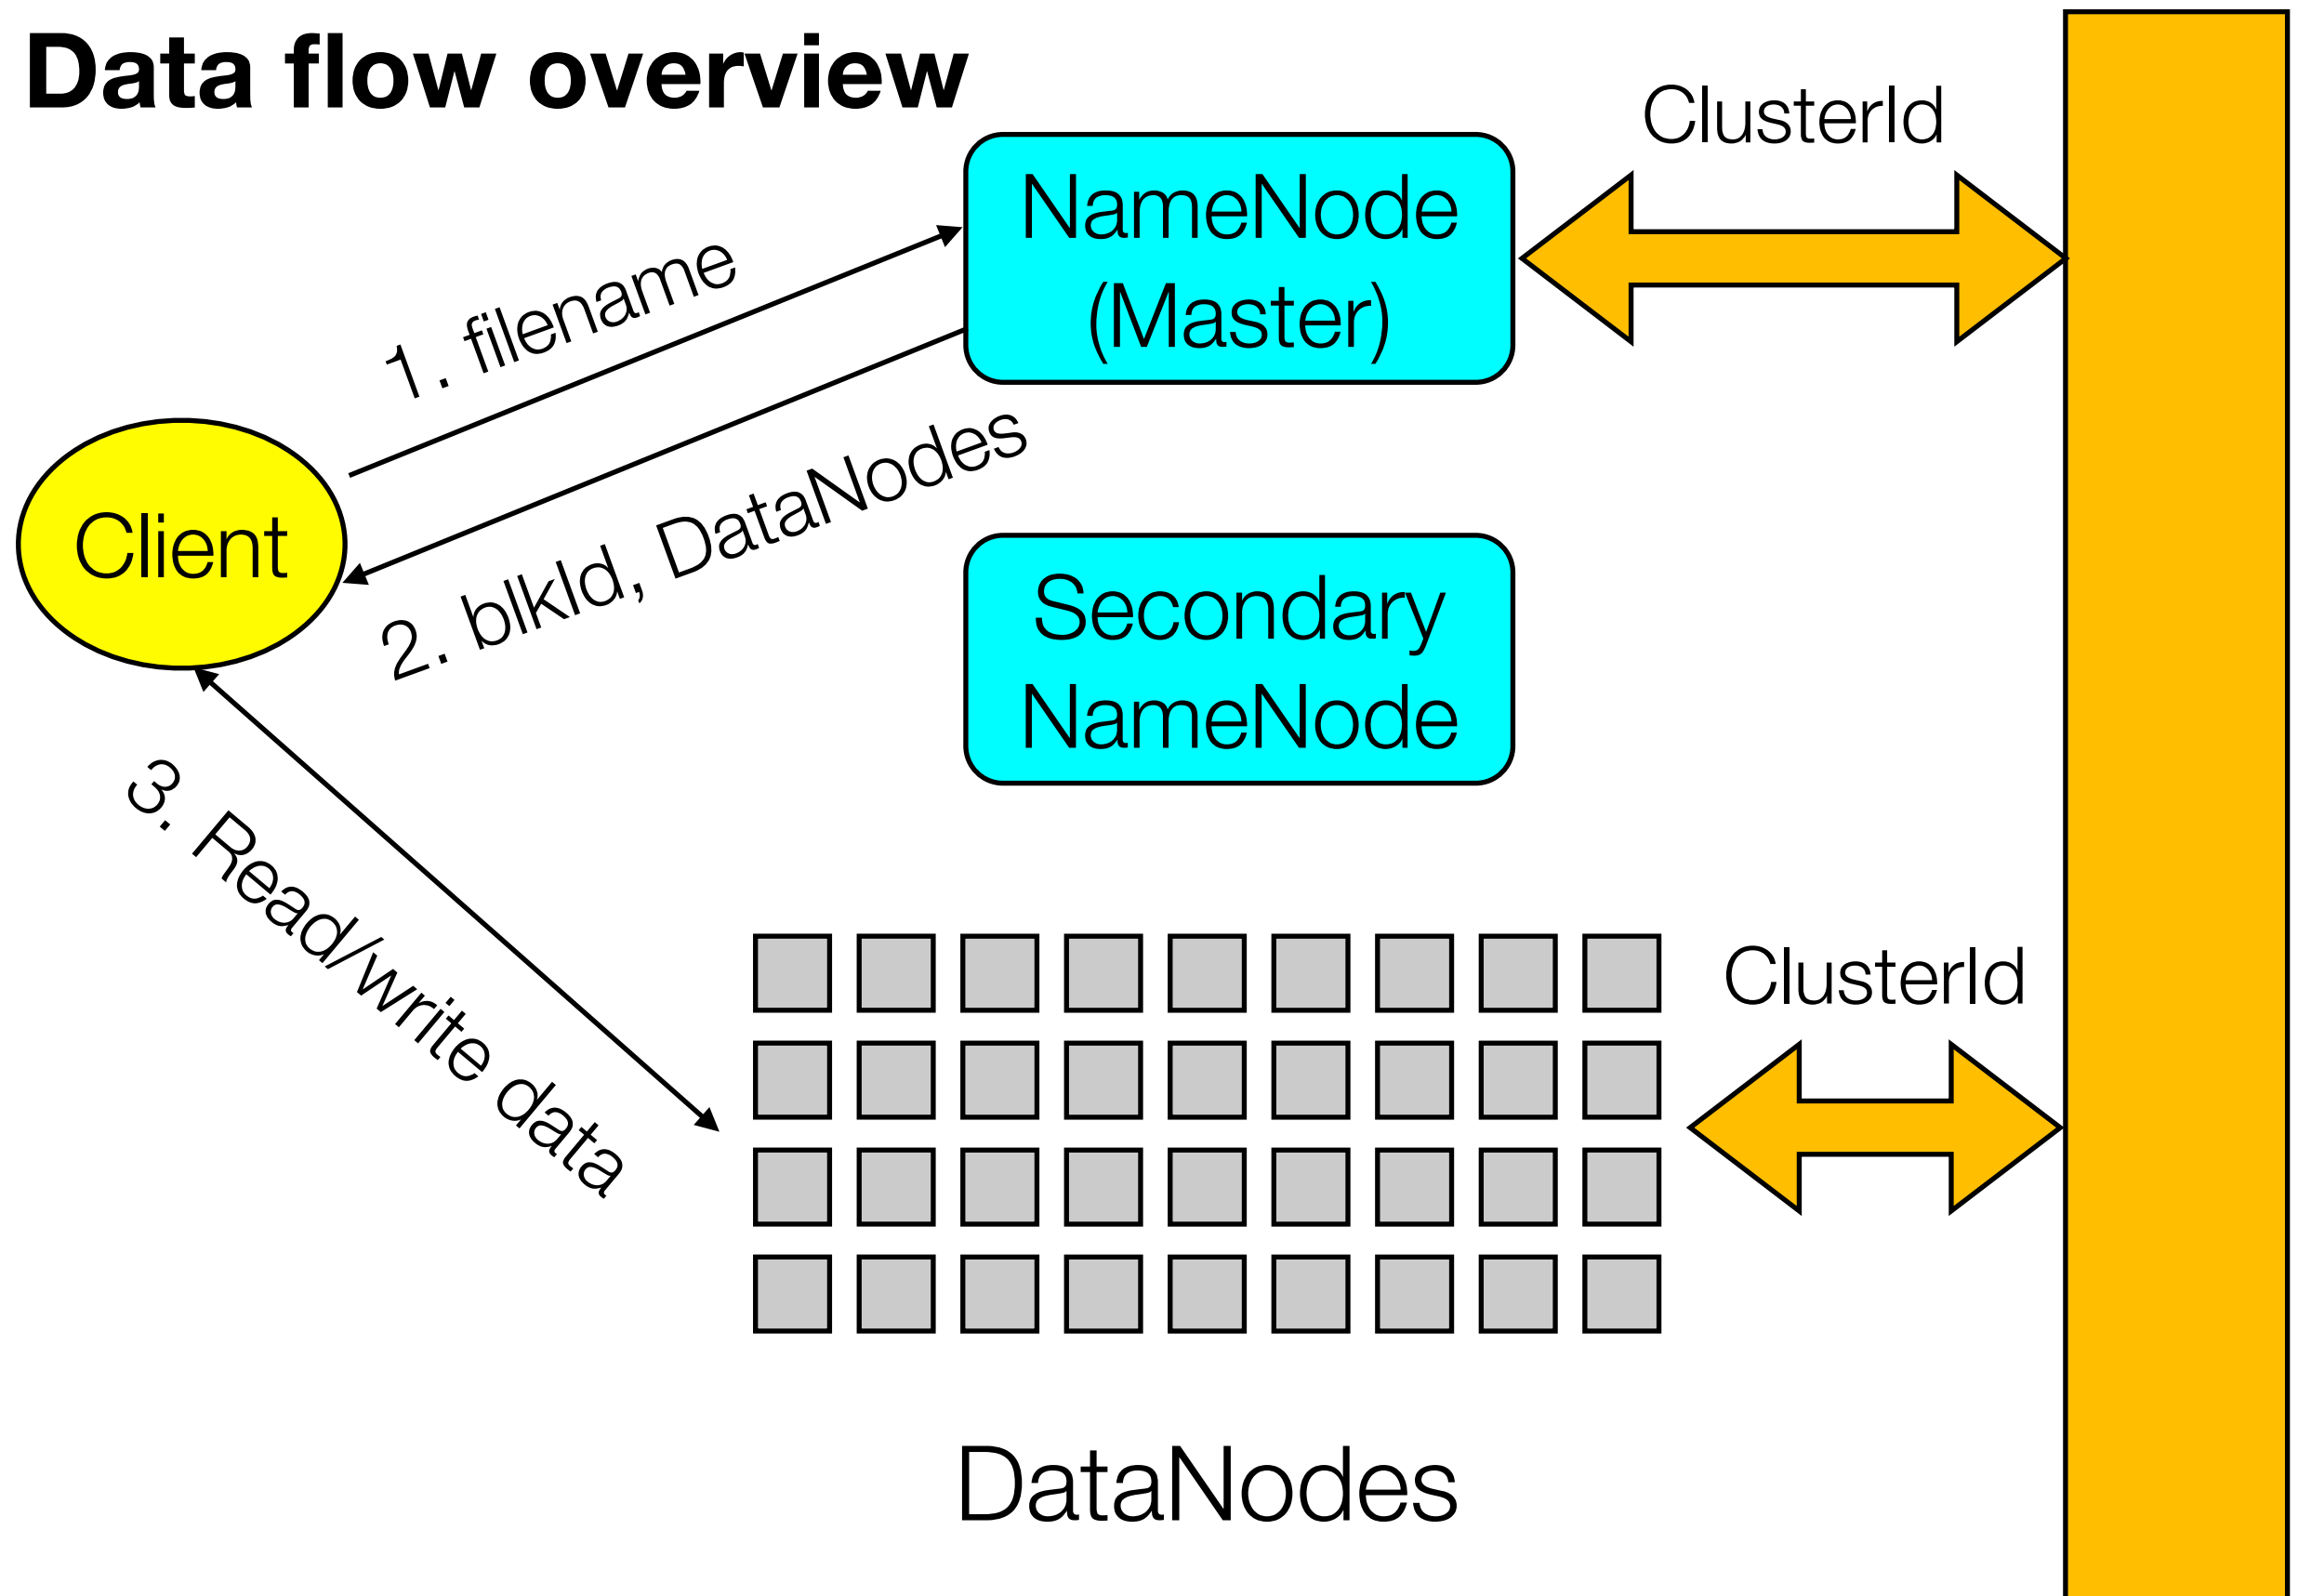
- ▶ Provides many interfaces and libraries for different filesystems
 - ▶ S3, KFS, etc.
 - ▶ Thrift (C++, Python, ...), libhdfs (C)

Architecture

HDFS Architecture



Data flow overview



NameNode (master)

Functions of a NameNode

Manages filesystem namespace

- ▶ Maps a filename to a set of blocks
- ▶ Maps a block to the DataNodes where it resides

Cluster configuration management

Replication engine for blocks

NameNode metadata

Metadata kept in memory

Types of metadata

Transaction log

- ▶ List of files
 - ▶ List of blocks for each file
 - ▶ File attributes, e.g., creation time, replication factor
-
- ▶ List of DataNode for each block

A transaction log: records file creations, file deletions, etc.

DataNode

DataNode

Stores data in the local file system

Stores metadata of a block (e.g., CRC checksum)

Serves data and metadata to clients

Communicates with NameNode through periodic “heartbeat” (once per 3 secs)

DataNode

Block report

- ▶ Periodically (1-hour by default) sends a report of all existing blocks to the NameNode

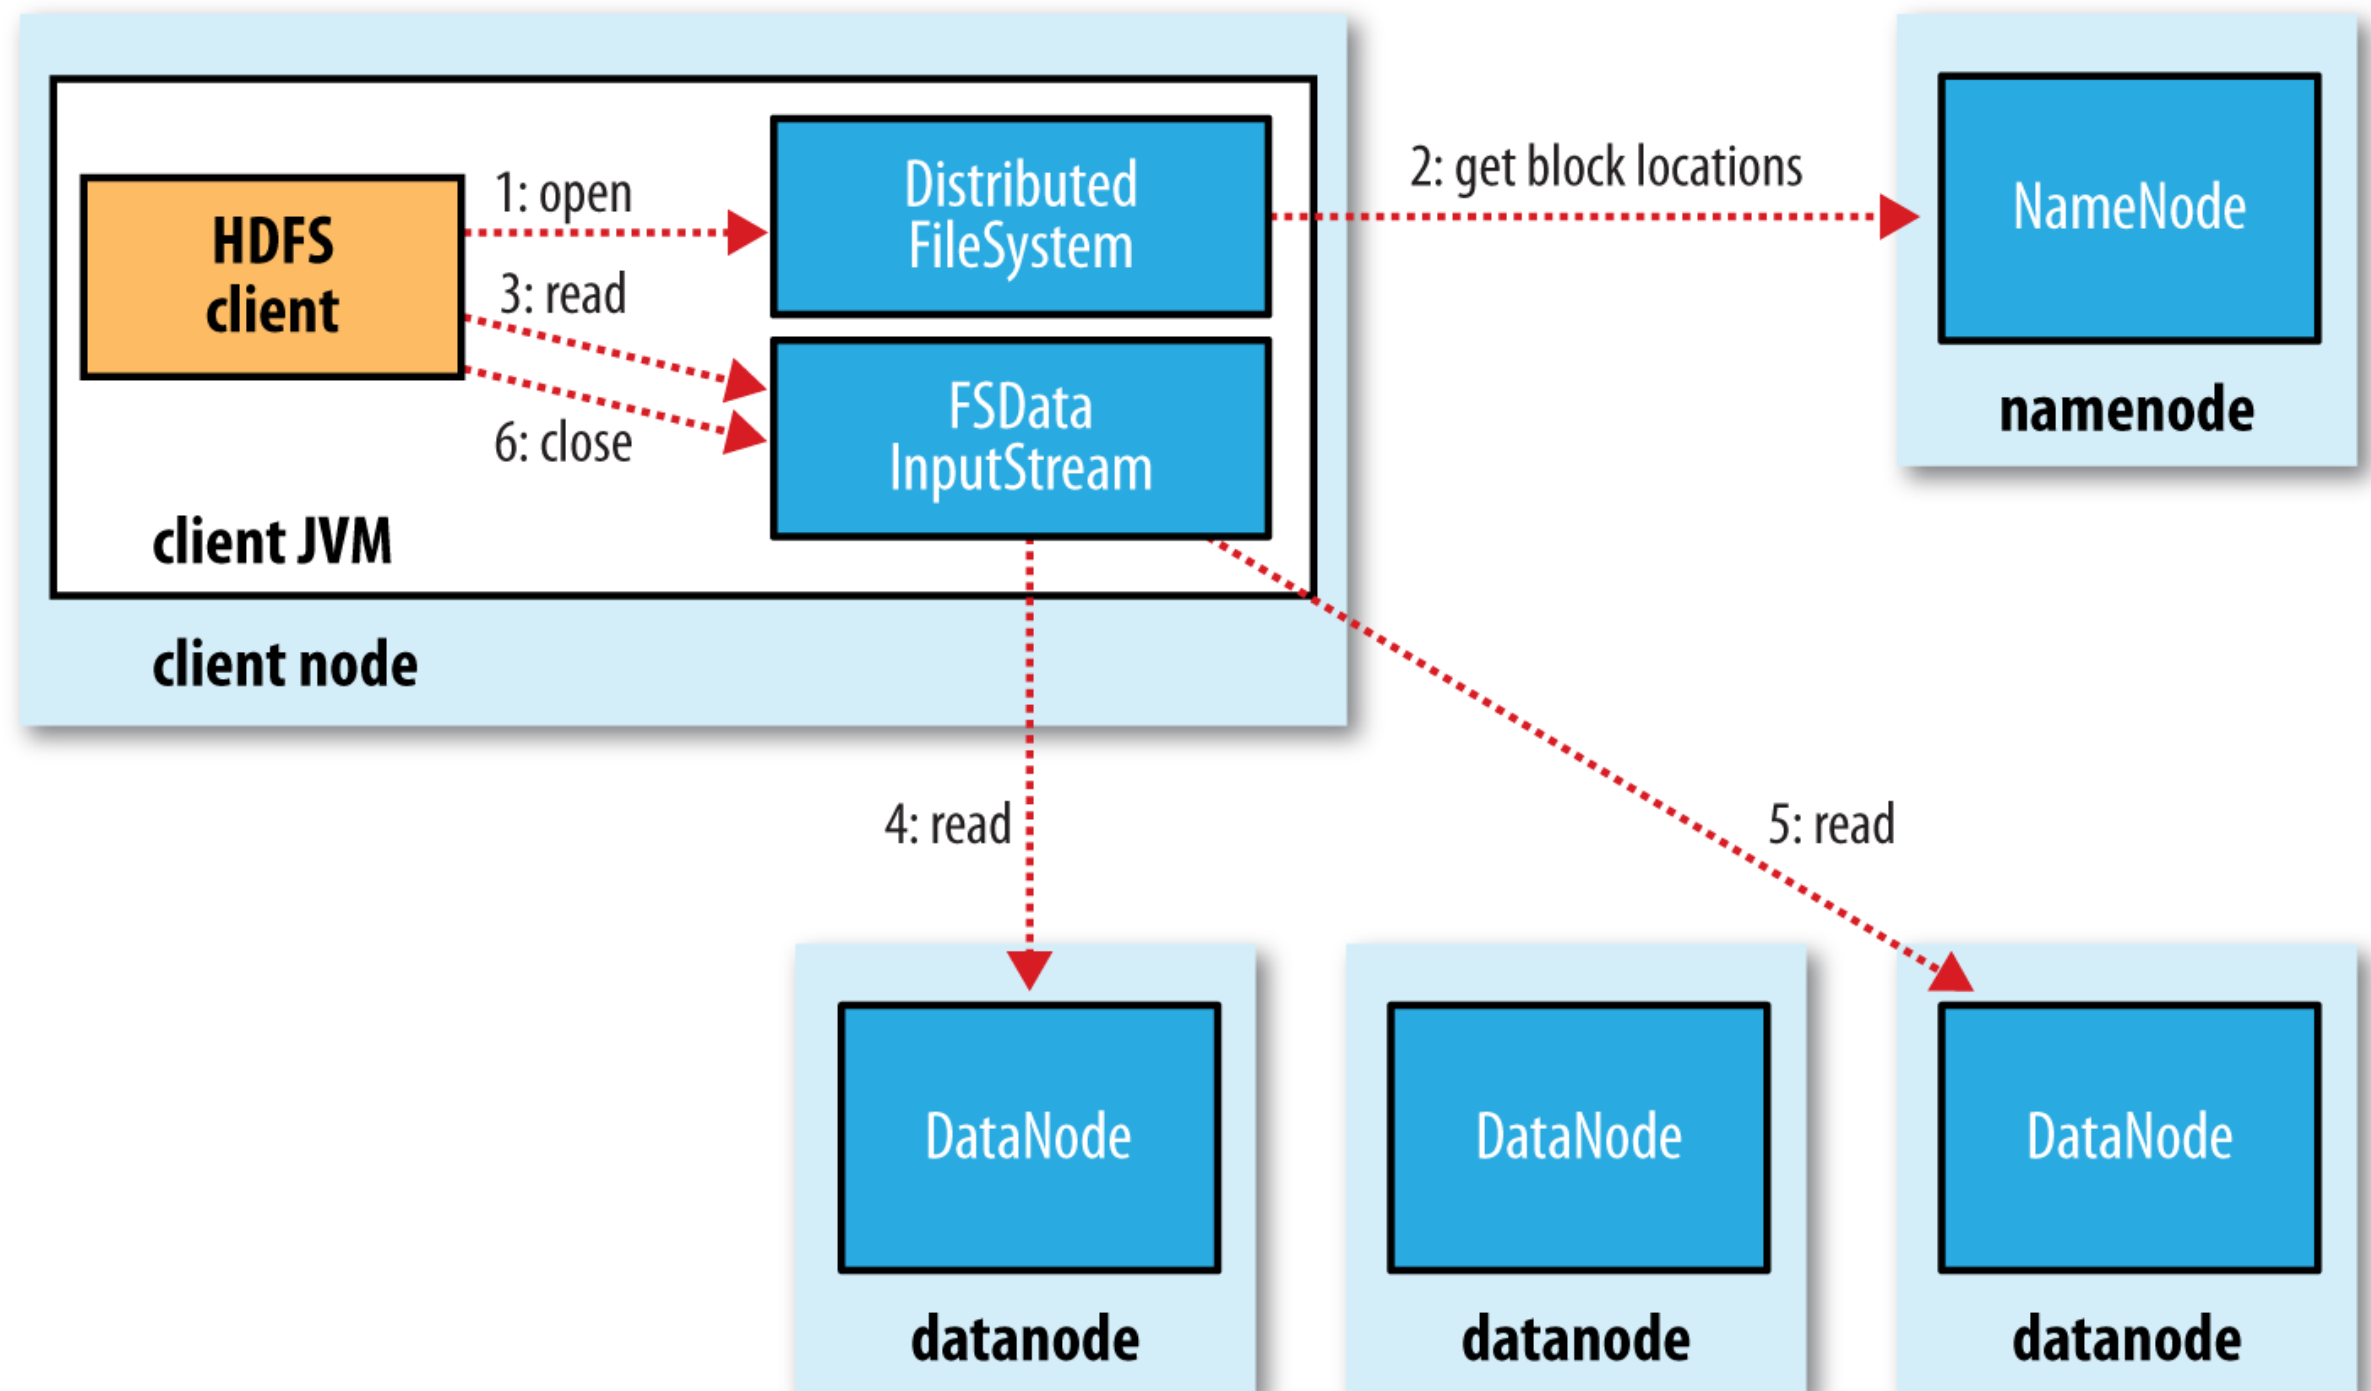
Facilitates pipelining of data

- ▶ Forwards data to other specified DataNodes

HDFS Workflow


File read and write

Anatomy of a file read



How to choose the “closest”
block?

Choosing the “closest” block



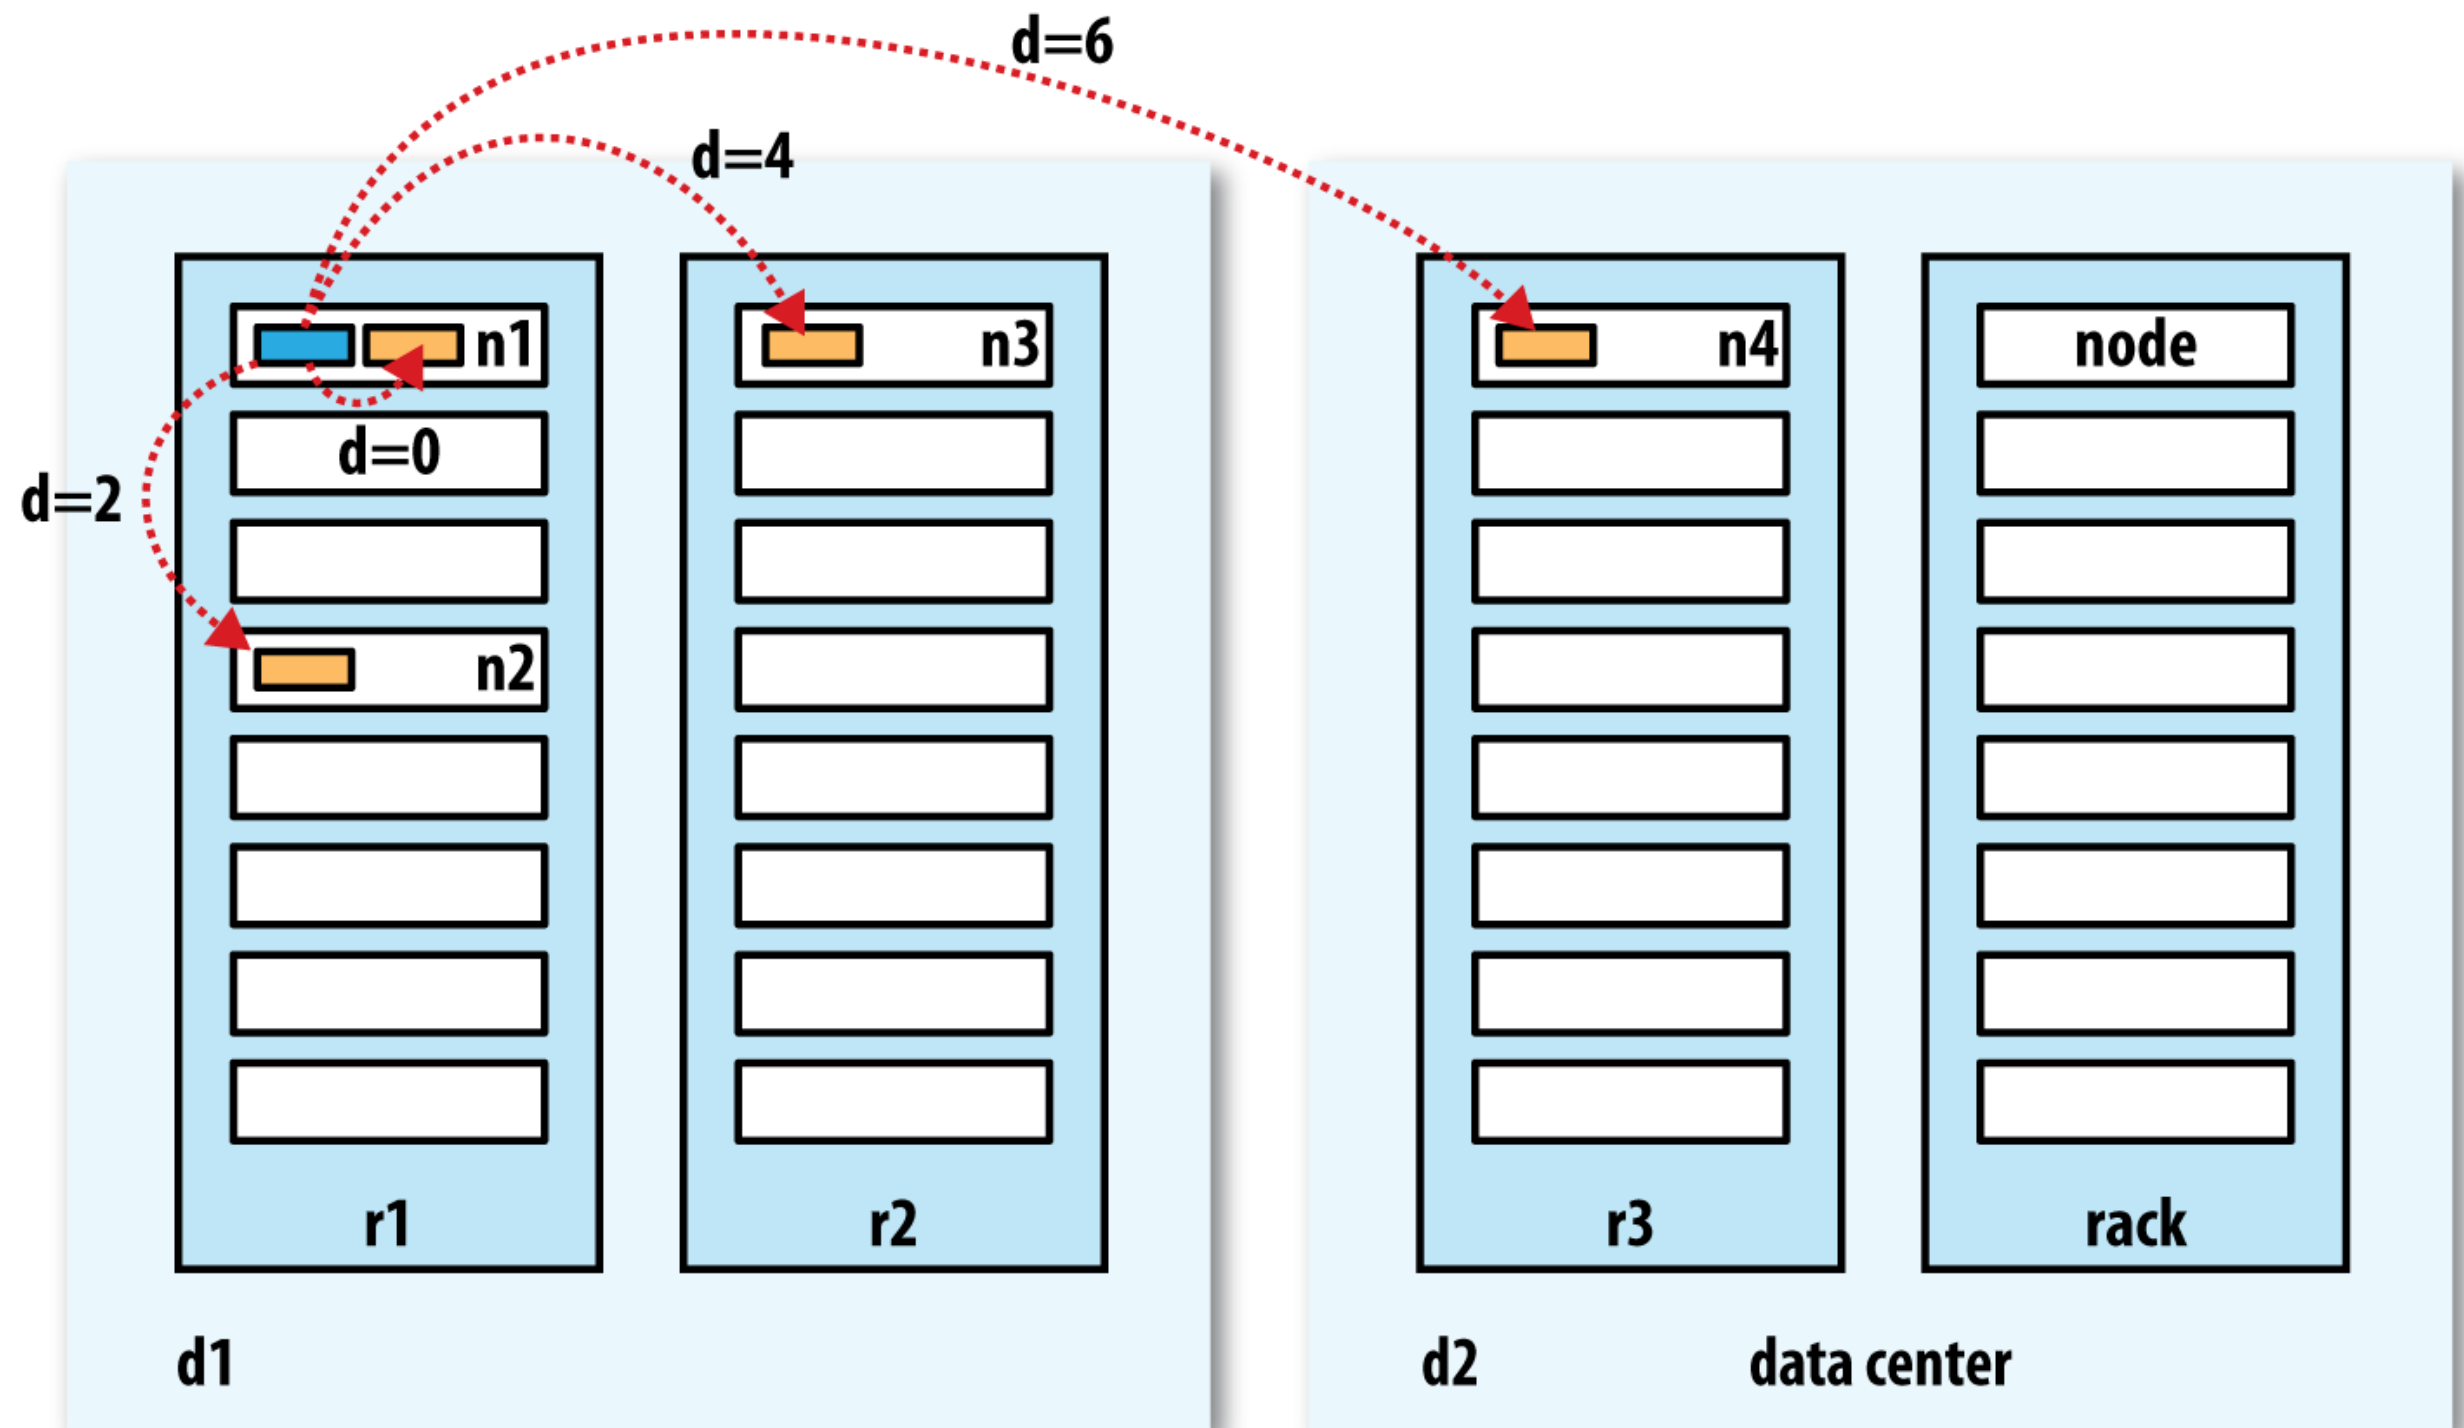
Choosing the “closest” block

Computing the ***distance*** between two nodes

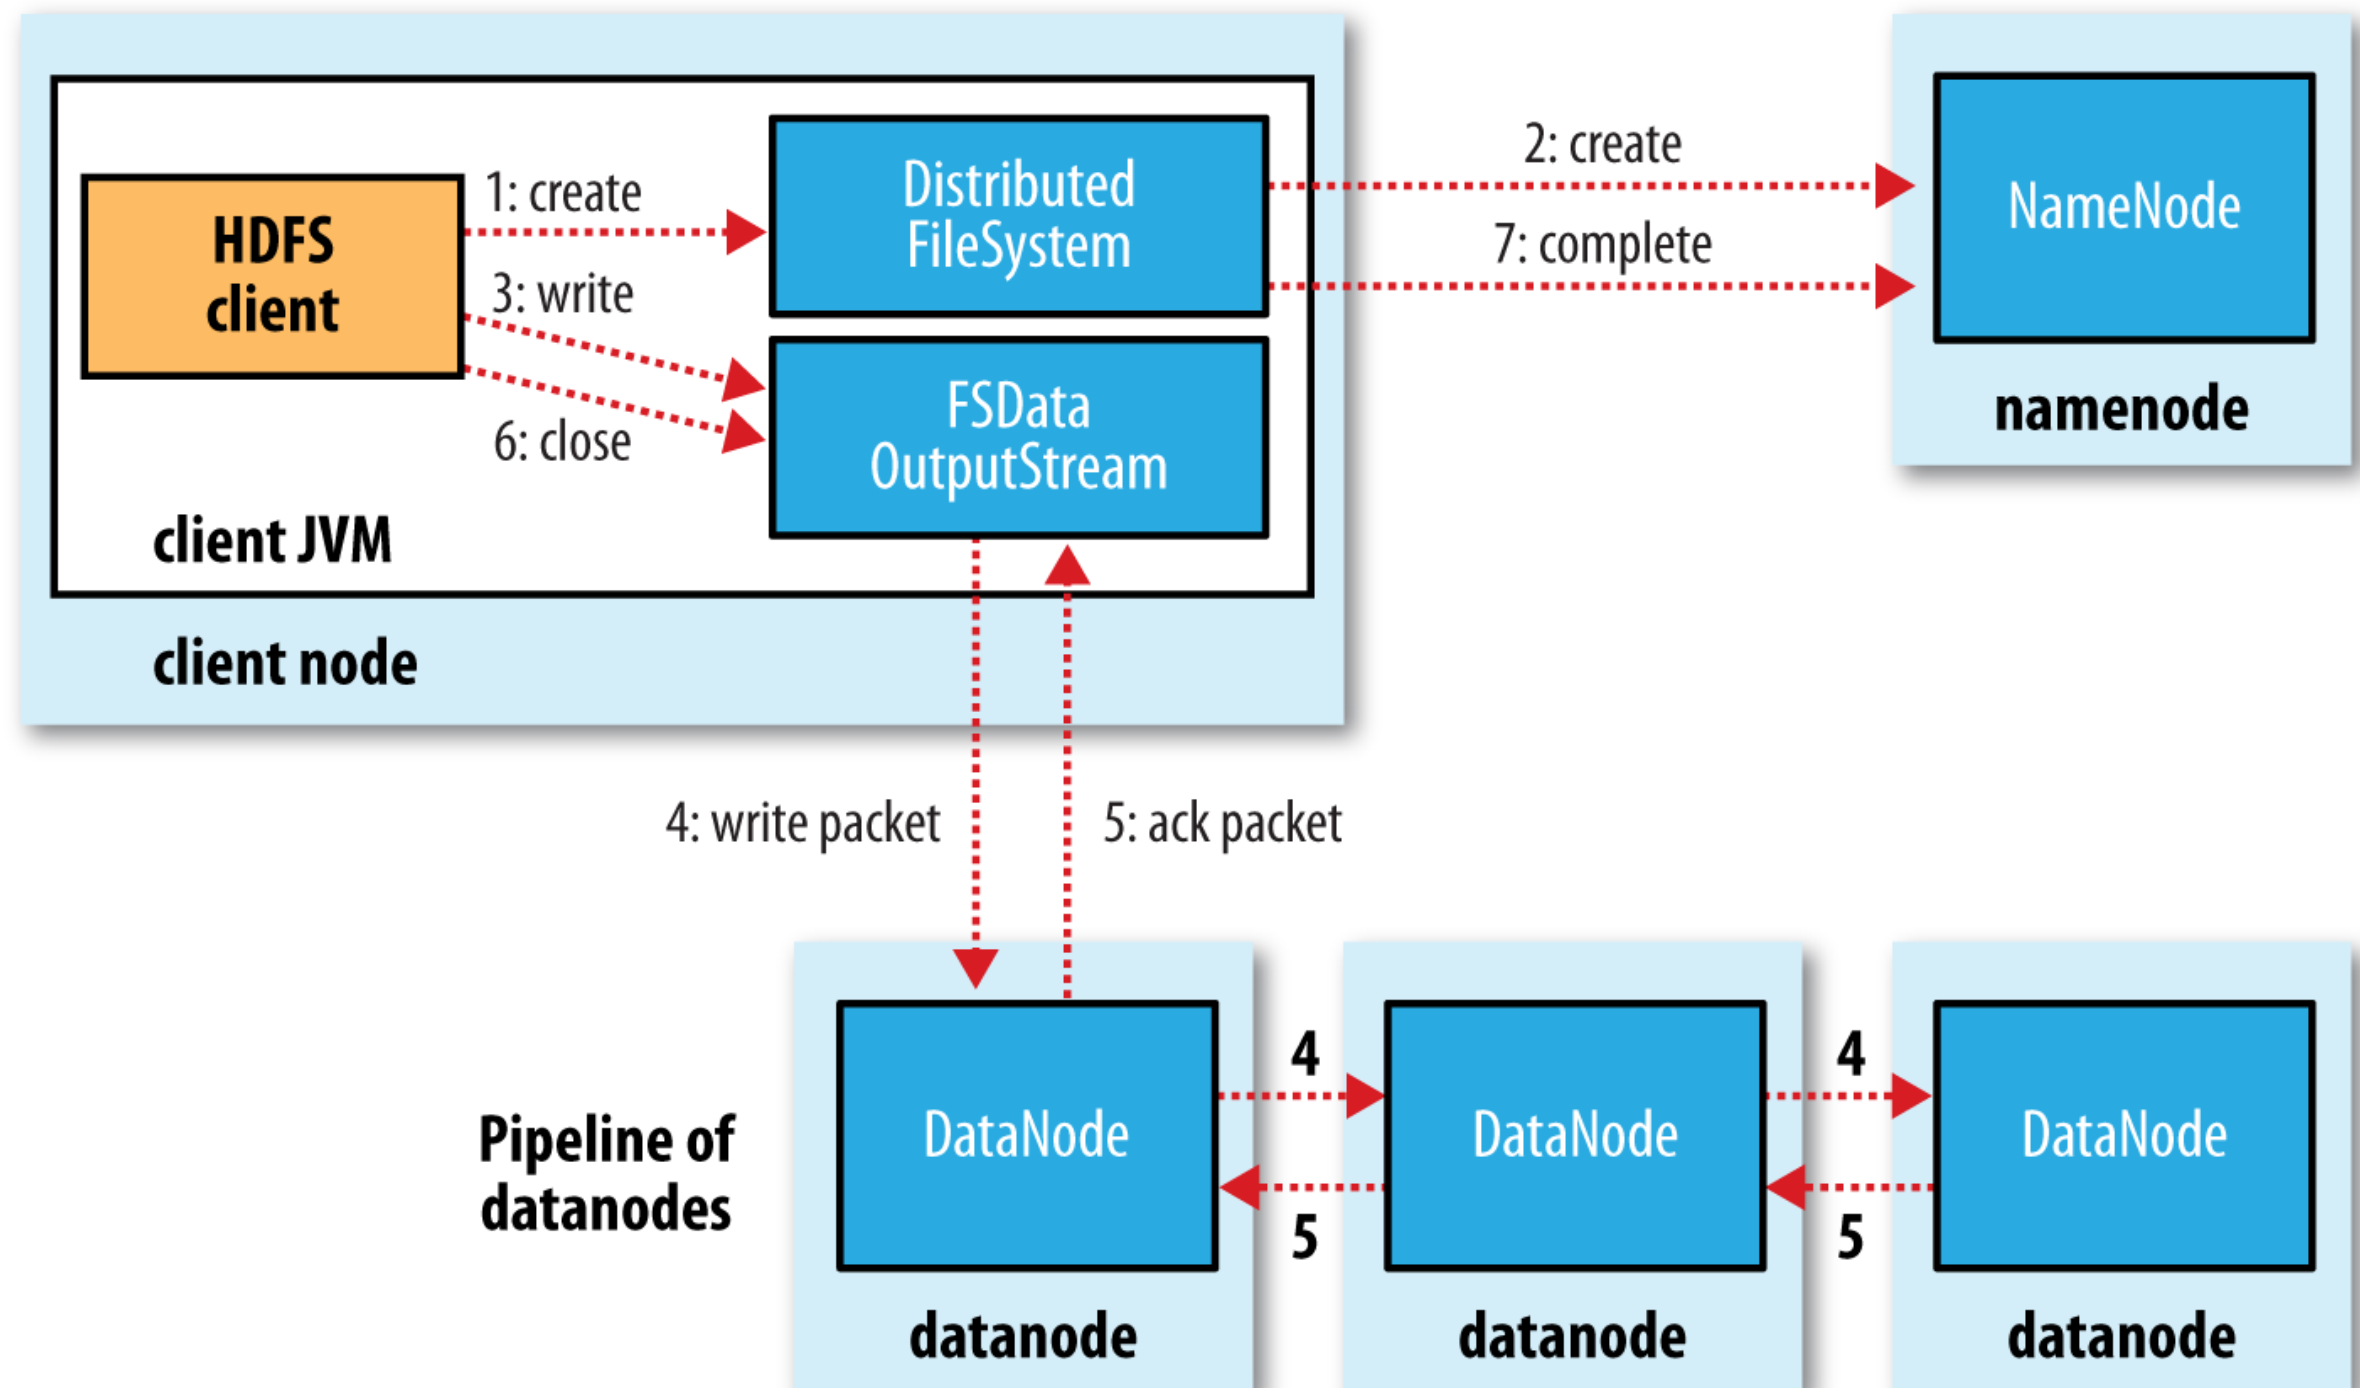
Denote a node $n1$ on rack $r1$ in DC $d1$ by $/d1/r1/n1$

- ▶ $\text{dist}(/d1/r1/n1, /d1/r1/n1) = 0$ (process on the same node)
- ▶ $\text{dist}(/d1/r1/n1, /d1/r1/n2) = 2$ (different nodes on the same rack)
- ▶ $\text{dist}(/d1/r1/n1, /d1/r2/n3) = 4$ (nodes on different racks in the same datacenter)
- ▶ $\text{dist}(/d1/r1/n1, /d2/r3/n4) = 6$ (nodes in different datacenters)

Distance between two nodes



Anatomy of a file write



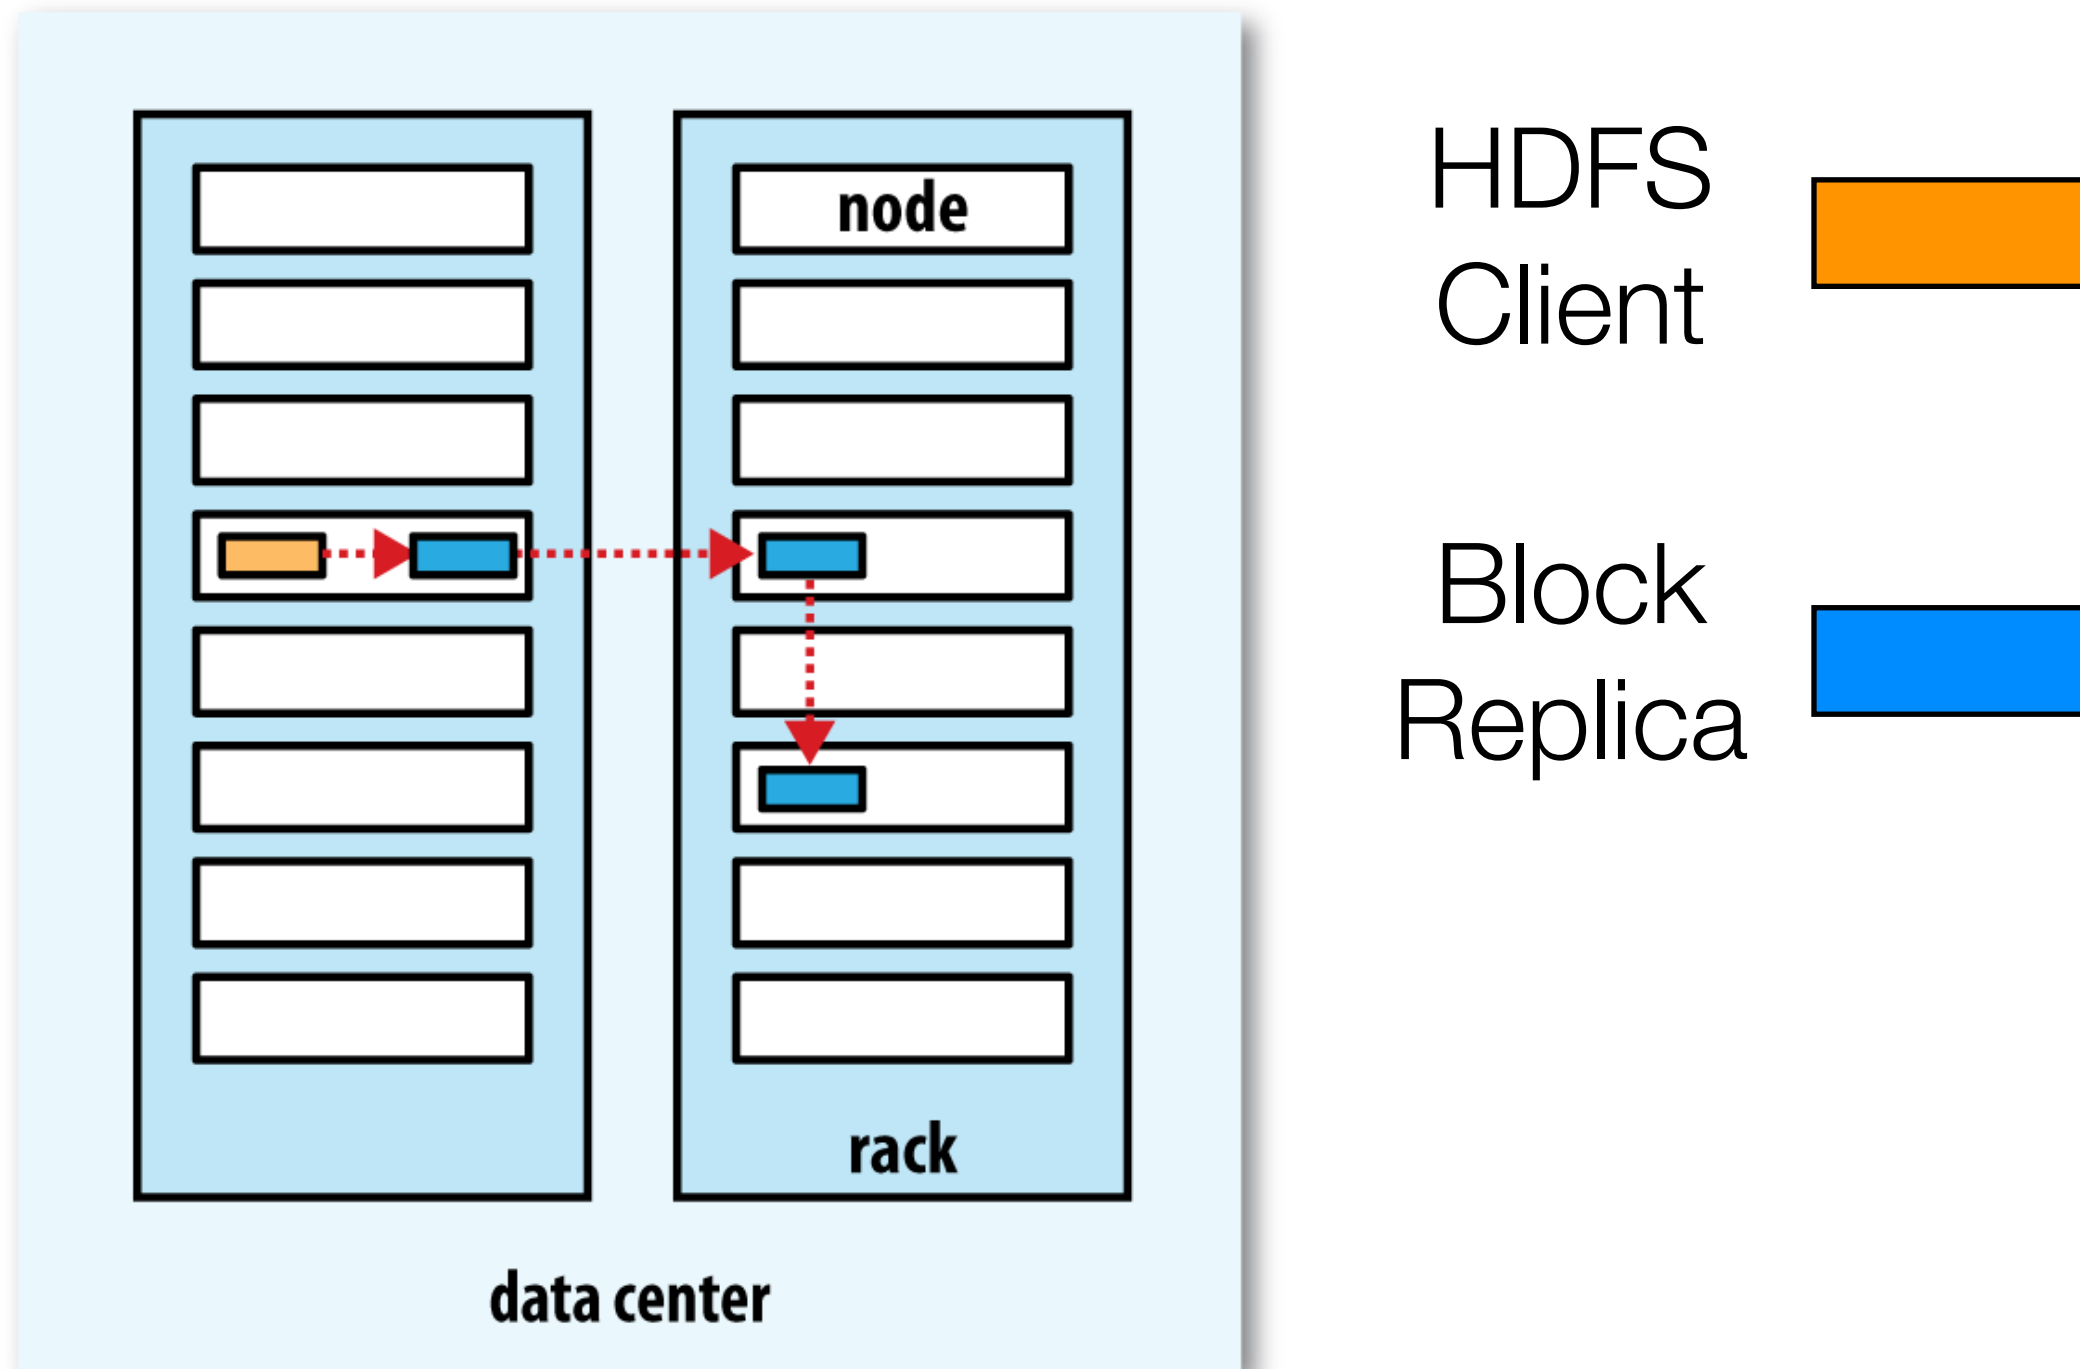
Block placement

Current strategy (replaceable with customized policy)

- ▶ One replica on local node
- ▶ 2nd and 3rd replica on two nodes of same *remote* rack
- ▶ Additional replicas are randomly placed

Once the replica locations
have been chosen, a
pipeline is built

Replica pipeline



Handle failures

Heartbeats

DataNodes send heartbeats to the NameNode

- ▶ Once every 3 secs

NameNode uses heartbeats to detect DataNode failure

- ▶ No response in 10 mins is considered a failure

Replication engine

Upon detecting a DataNode failure

- ▶ Choose new DataNodes for replicas
- ▶ Balance disk usage
- ▶ Balance communication traffic to DataNodes

Data corrections

Checksums to validate data (CRC32)

File creation

- ▶ Client computes checksum per 512 byte
- ▶ DataNode stores the checksum

File access

- ▶ Client retrieves data and checksum from DataNodes
- ▶ If validation fails, try other replicas

NameNode failure

A single point of failure

Transaction log stored in multiple directories

- ▶ Directory on local file system
- ▶ A directory on a remote file system (NFS)

Add a **secondary NameNode**

Secondary NameNode

Not Standby/Backup NameNode

- ▶ only for checkpointing
- ▶ has a NON-Realtime copy of FSImage

Copies NameNode's FSImage & Transaction Log

Merges them to a new FSImage

Uploads new FSImage to the NameNode and purges Transaction Log

Summary

As an open-source implementation of GFS, HDFS shares the same design assumptions

- ▶ Very large files
- ▶ Streaming data access pattern
- ▶ Commodity hardware

Limitations

Low-latency data access

- ▶ tens of millisecond range
- ▶ HDFS emphasizes throughput over latency

Lots of small files

- ▶ billions of files
- ▶ All meta data are kept in memory, resulting in overflow

Multiple writers, arbitrary file modifications

HDFS FileSystem API (Java):
`org.apache.hadoop.fs`

Reading data

A general filesystem API is provided by `FileSystem`

Retrieve an instance using static factory methods:

```
public static FileSystem get(URI uri, Configuration conf) throws IOException
```

used to infer the filesystem scheme (e.g.,
`hdfs://` for HDFS, `file:///` for local filesystem)



encapsulates a client's config, usually
set in `etc/hadoop/core-site.xml`

Reading data

A file in HDFS is represented by a Hadoop `Path` object

- ▶ HDFS URI: `hdfs://localhost/user/weiwa/hkust.txt`

Get the input stream of a file using `open()` method


```
public FSDataInputStream open(Path f) throws IOException
```

```
public abstract FSDataInputStream open(Path f, int bufferSize)  
throws IOException
```

Reading data

Putting them together: `cat` a file

```
public class FileSystemCat {  
  
    public static void main(String[] args) throws Exception {  
        String uri = args[0];  
        Configuration conf = new Configuration();  
        FileSystem fs = FileSystem.get(URI.create(uri), conf);  
        InputStream in = null;  
        try {  
            in = fs.open(new Path(uri));  
            IOUtils.copyBytes(in, System.out, 4096, false);  
        } finally {  
            IOUtils.closeStream(in);  
        }  
    }  
}
```



A handy I/O tool

Writing data

Use the create() method

public FSDataOutputStream **create**(Path f) **throws** IOException

```
String localSrc = args[0];
```

```
String dst = args[1];
```

```
Configuration conf = new Configuration();
```

```
LocalFileSystem localFS = LocalFileSystem.get(conf);
```

```
FSDataInputStream in = localFS.open(new Path(localSrc));
```

```
FileSystem outFS = FileSystem.get(URI.create(dst), conf);
```

```
FSDataOutputStream out = outFS.create(new Path(dst));
```

```
IOUtils.copyBytes(in, out, 4096, true);
```

Deleting data

Use the `delete()` method on `FileSystem` to permanently remove files or directories:

public boolean delete(Path f, **boolean** recursive) **throws** IOException

- ▶ recursive is ignored if f is a file or an empty directory
- ▶ returns true if delete is successful

```
String uri = args[0];  
Configuration conf = new Configuration();  
FileSystem fs = FileSystem.get(URI.create(uri), conf);  
fs.delete(new Path(uri), true);
```

Credits

Some slides are adapted from Dhruba Borthakur's slides