

CSIT6000P Spatial and Multimedia Databases
2022 Spring



香港科技大學
THE HONG KONG UNIVERSITY OF
SCIENCE AND TECHNOLOGY

Spatial Data Organization & Indexing

Prof Xiaofang Zhou

+ Learning Objectives

■ What we will cover

- Basic principles of managing multidimensional data
- Representative data access methods for point and polygon data
- Processing of some simple spatial operations using data access methods

■ Goals

- Understand major types of multidimensional data access methods, and their strengths and limits
- Understand how point query, window query and join query are processed using various data access methods

+ Readings

- R. Güting, An Introduction to Spatial Database Systems, *The VLDB Journal*, 3:4, 1994
- V. Gaede and O Günther, Multidimensional Access Methods, *ACM Computing Surveys*, 30:2, 1998
- J. Orenstein and F. Manola, PROBE Spatial Data Modeling and Query Processing in an Image Database Application, *IEEE Transactions on Software Engineering*, 14:5, 1988
- T. Brinkhoff, H.-P. Kriegel and B. Seeger, Efficient Processing of Spatial Joins Using R-Trees, SIGMOD 1993

+ Spatial Indexing

- Purpose:

- **Efficiency** in processing spatial selection, join and other spatial operations

- Two strategies to organize space and objects


- Dedicated external data structures
- Map spatial objects into 1D space and use a standard index structure in RDBMS (e.g., B⁺-tree)

- Basic ideas

- **Approximation** and **hierarchical** data organisation

+ Object Approximation

- A fundamental idea of spatial indexing is the use of approximation
- **Object** approximation
 - Object centric
 - EG: use of **MBRs** (minimum bounding rectangles)
- **Grid** approximation
 - Space centric
 - EG: **quadtree**



+ What Do We Need to Know?

For each multidimensional access method:

- Motivation

- Why is it proposed?
- Good for what type of data?
 - points? polygons?

- Operations for **creating** and **maintaining** an index

- Insert and delete data items

- Query operations using an index

- Point query and window query
- Spatial join queries and other queries

+ Background Knowledge

- Disks and files
- Basic indexing methods in relational DBMS
 - B⁺-Tree
 - Hashing
 - Bitmap
- Query processing using indexes
 - What to achieve, and how?

+ Background: Disks and Files

- (S)DBMS stores information on **hard disks**
- This has major implications for DBMS design:
 - **READ:** transfer data from disk to main memory
 - **WRITE:** transfer data from RAM to disk
 - Both are high-cost operations (I/O), so must be planned carefully!

+ Magnetic Disks


- Secondary storage device of choice
- Supports random data access (vs. sequential access such as using tapes)
- Data is stored and retrieved in units called disk **blocks** or **pages** or **buckets**
- Unlike RAM, time to retrieve a page varies depending upon their locations on the disk

... technologies have evolved significantly, but these principles still hold

+ Accessing a Disk Block

10

- Time to access (read/write) a disk block:
 - **seek time**: moving arms to position disk head on the track
 - **rotational delay**: waiting for the block to rotate under the head
 - **transfer time**: moving data to/from the disk surface of the block
- Seek time and rotational delay dominate
- Key to lower I/O cost: reduce seek/rotation delays!
 - Hardware vs. software solutions?



+ Index Structure

- An index is an **auxiliary** file that makes it more efficient to search for a record in the data file
 - Think about the index at the end of a book
- An index is usually specified on one field of the file (although it could be specified on several fields)
 - It's called the search key, or simply **key**
- One form of an index is a file of entries **ordered** by key values: `<key, pointer >`
 - The pointer can point to a physical or logical address (of a page)

+ Costs and Benefits of Using Indexes

12

- An index always has two costs
 - Space to store it
 - The time to create and maintain it
- An index file typically occupies less disk blocks than the data file
 - Because its entries are much smaller
 - Therefore, index files are often kept in memory
- A binary search (for the sorted indexes) on the index yields a pointer to the file record
 - Then, an extra I/O access is needed to fetch the data

+ Data Access Methods

13

■ One dimensional

- Hashing and B⁺-Trees

■ Point data

- Hashing: GRID and EXCELL
- Hierarchical
 - Quadtree: point and region quadtrees
 - kd-Tree
 - Z-values and B+-tree

■ Polygon data

- Transformation: end point mapping and z-values
- Overlapping: R-tree and R*-tree
- Clipping: R+-tree

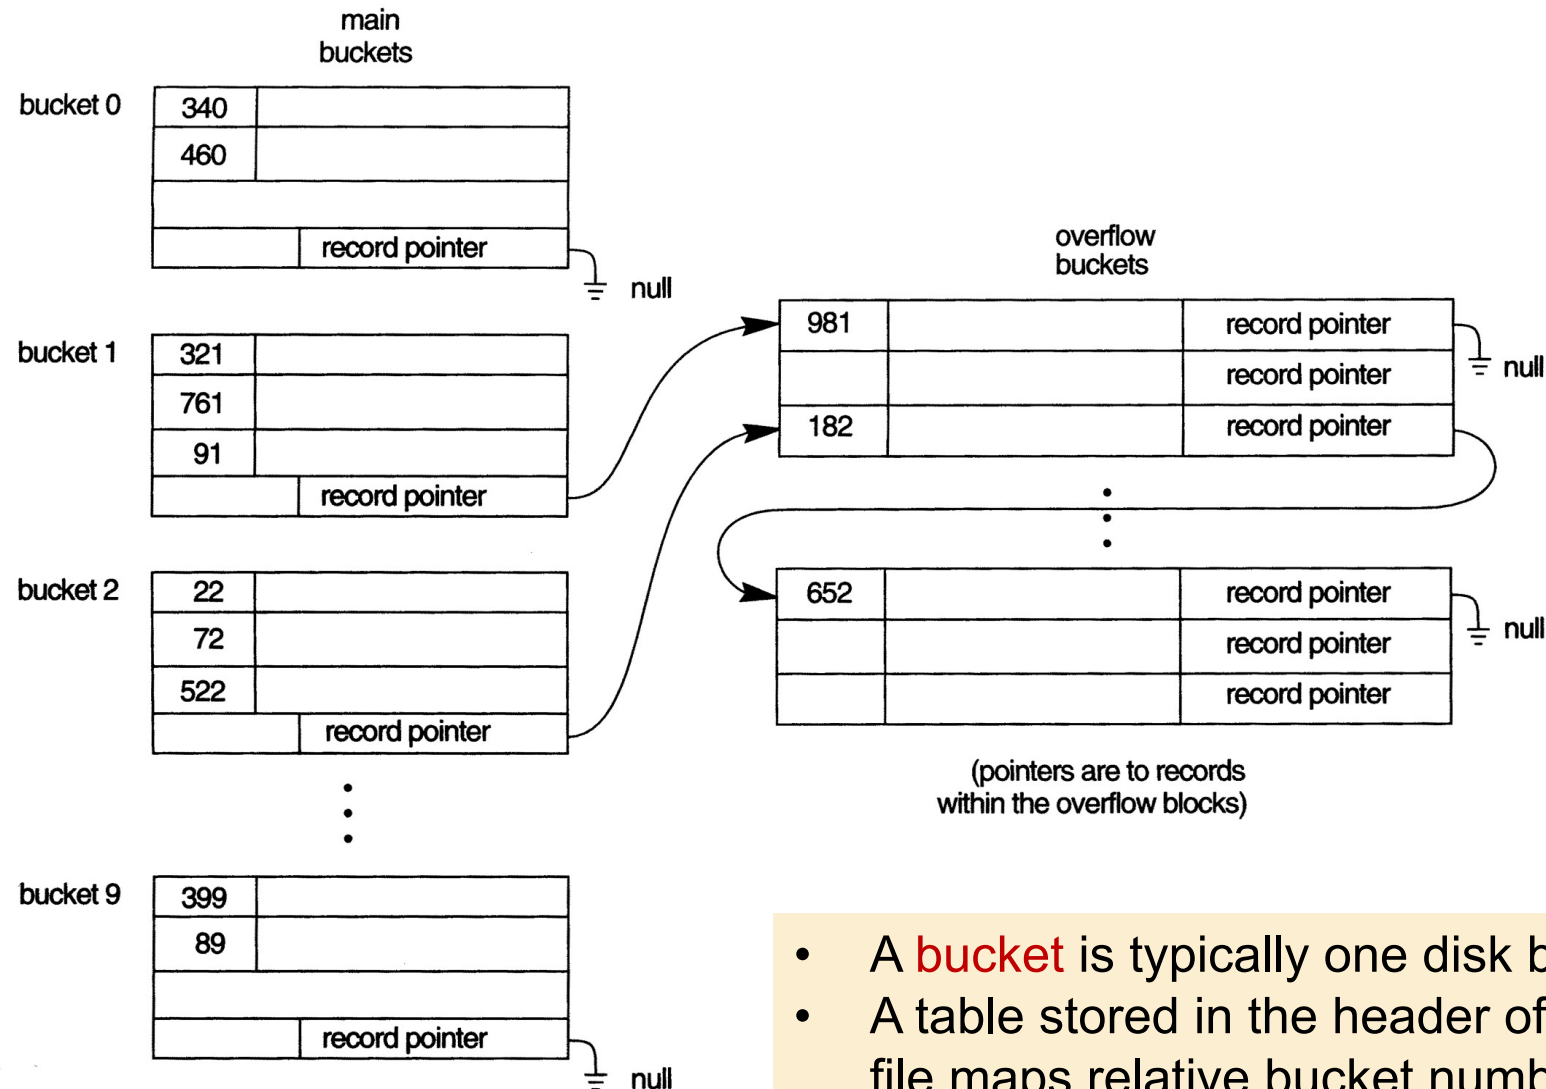
+ Static Hashing

14

- Hashing converts the key of a record to an address in which the record is stored
 - A **hash function** is used to map the **key** to the relative address of a bucket in which the record is stored
- If a file is allocated with m buckets, the hash function must convert a key k into the relative address of the block:
 - $h(k)$ in $\{0, \dots, m-1\}$
 - For example, for integer key values, $h(k) = k \bmod m$
- Search cost?
 - Can be $O(1)$

+ A Hashing Index Example

15




- A **bucket** is typically one disk block
- A table stored in the header of the hash file maps relative bucket numbers to actual disk block addresses

+ Collision

16

- Insertion of a new record may lead to collision
 - No space in $b = h(k)$
- **Chaining:** Use overflow buckets
 - Example: Common overflow area for all blocks in a file
 - Each block has a pointer to its first record in the overflow area
 - Records belonging to the same block are linked by pointers



- What problems may this cause?
- Why does this happen?

+ Hashing Functions

17

- A good hash functions must
 - be computed efficiently
 - minimize the number of collisions by spreading keys around the file as uniform as possible
- Example of hash functions
 - **Truncation**: take the first k bits of a key, for 2^k buckets
 - Taking the first 2 bits, there are $2^2 = 4$ buckets: 00, 01, 10, 11
 - **Division**: $h(k) = k \bmod m$, where m is the number of buckets

... for the truncation method, one can also takes the last k bits. Compare and contrast these two truncation methods

+ Pros and Cons of Hashing

18

- Excellent performance for searching on **equality** on the key used for hashing
 - Records are not ordered, thus any search other than on equality is very expensive
- Choose a good hash function is difficult
 - Must assume the data distribution (now and in the future)
 - Prediction of total number of buckets is difficult
 - Too few or too many are both bad
 - A practical approach is to estimate a “reasonable” size and periodically reorganize
- Recently strong renewed interest in hashing
 - Learn-to-hash, order-preserving hashing, and locality-sensitive hashing (LSH)

+ Linear Hashing

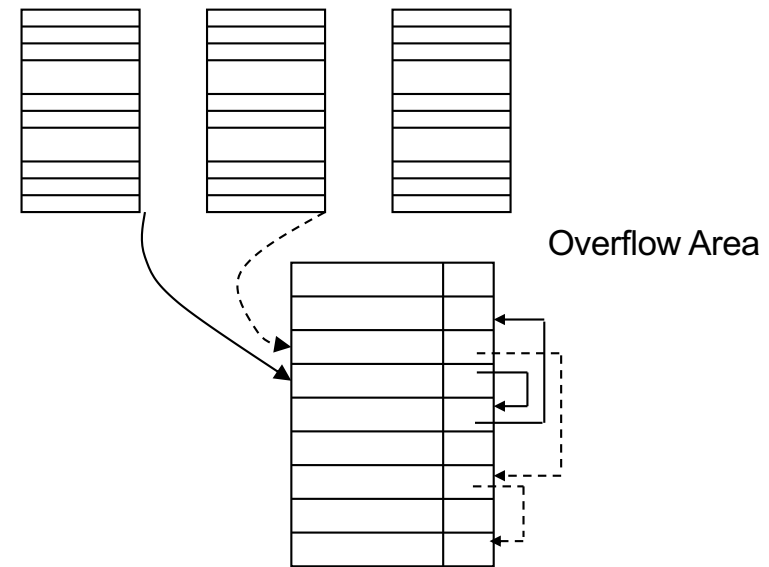
- The file size grows linearly, bucket by bucket
 - The file starts with m main buckets, numbered from 0 to $m-1$
- Initial hashing function $h_1(k) = k \bmod m$
- We keep the following info:
 - n : a pointer to the bucket that should be split next
 - Initially, $n = 0$

+ Insertion of a Record

20

If there is a collision, do:

- Push the record to an overflow bucket
- (**Grow**) A new bucket is appended at the end of the hash table
- (**Split**) Records in bucket number n (including those in the overflow space) are hashed again using
 - $h_2(k) = k \bmod (2m)$ (this will return either n or $m+n$, why?)
 - That is, these records will either remain in bucket n or in bucket $m+n$
- If $n = m$ (all original buckets have been split), we will set:
 - $n = 0$
 - $h_1(k) = k \bmod (2m)$
 - $h_2(k) = k \bmod (4m)$



For all k values such that $h_1(k) = k \bmod m = n$, they are in the form of $k = i \times m + n$.

Now $h_2(k) = k \bmod (2m)$

When i is an odd number, $h_2(k) = m+n$


When i is an even number, $h_2(k) = n$

+ Search for a Record

- Search for a record with field value k
 - $l = h_1(k)$
 - if $l < n$ then $l = h_2(k)$ (i.e., bucket l has been split already)
 - search the bucket whose hash value is l
- Advantage
 - The file size grows linearly, bucket by bucket
- Disadvantage
 - No guarantee that a split relieves the overloaded bucket
 - A split is not at where the bucket currently overflows, but at the “next” bucket in a round-robin fashion
 - Still requires overflow buckets and chaining


+ Extendible Hashing

- The file is structured into two levels:
 - **directory** and **buckets**
- The directory has 2^d entries (d : **global depth**)
 - Each entry points to a bucket
 - Use some hashing function that generates a string of bits
 - The first d digits used as index into the directory
 - Several entries can point to the same bucket
- **local depth** d' ($d' \leq d$)
 - Only the first d' bits are used



+ Extendible Hashing Adjustment

- The directory size can be doubled or halved
 - Double: $d = d + 1$, when a bucket with $d' = d$ overflows
 - Half: $d = d - 1$ when $d > d'$ for all buckets
- Two-level search, highly efficient
 - No just-in-case space waste, no chain search



... practically all modern file systems use either extendible hashing (Ronald Fagin, 1979) or B-trees

+ Data Access Methods

24

- One dimensional

- Hashing and B-Trees

- Point data

- Hashing: GRID and EXCELL
 - Hierarchical
 - Quadtree: point and region quadtrees
 - kd-Tree
 - Z-values and B-tree

- Polygon data

- Transformation: end point mapping and z-values
 - Overlapping: R-tree and R*-tree
 - Clipping: R+-tree

+ Grid File


25

■ Basic idea

- Superimpose a k -dimensional grid on the space
- Cells can be of variable sizes
- Cell-to-bucket mapping: **many-to-1**
 - *What about 1-to-many?*
- The **grid definition** is kept in memory
- The **grid directory** is kept on disk

■ Motivation

- A grid structure is simple, but a fixed-size grid is not suitable for non-uniformly distributed data



... what information to remember about the grid structure?

+ Grid File: Search

■ To answer a point query:

1. Use the **grid definition** to locate the cell
2. Find the reference information (i.e., a pointer to the data bucket)
3. Read the data bucket
4. Search in the data block

■ To answer a range query:

- Examine all the cells that overlap with the search region
- Read the corresponding cells to locate the data buckets(s)
- Read the data buckets, and search in these buckets

+ Grid File: Update

- To insert a point:

- Search (using a point query) the matching cell and the corresponding data bucket
- If there is sufficient space, insert into the data bucket
- If the data bucket is full, add a vertical or horizontal line to split cells, and redistribute data accordingly

- To delete a point:

- Can you do this?

+ EXCELL


28

■ Motivation

- Fixed grid is easier to manage and more efficient to use

■ Basic idea

- All cells are of the same size



Data buckets

... what information to remember about the grid structure?

+ EXCELL

- How to search for a point?
- How to search for all points in a given query window?
- How to insert a point?
- How to delete a point?

+ Data Access Methods

30

- One dimensional

- Hashing and B-Trees

- Point data

- Hashing: GRID and EXCELL
 - Hierarchical
 - Quadtree: point and region quadtrees
 - kd-Tree
 - Z-values and B-tree

- Polygon data


- Transformation: end point mapping and z-values
 - Overlapping: R-tree and R*-tree
 - Clipping: R+-tree

+ Uniform Decomposition

31

Recursive decomposition of space

Resolution: max. level of decomposition, leading to $2^n \times 2^n$ cells



An object can be represented as a collection of shaded cells

To have 1 x 1 cm cells, what is required resolution?

- an area of 5000 x 5000 km²
- an area of 300 x 300 km²

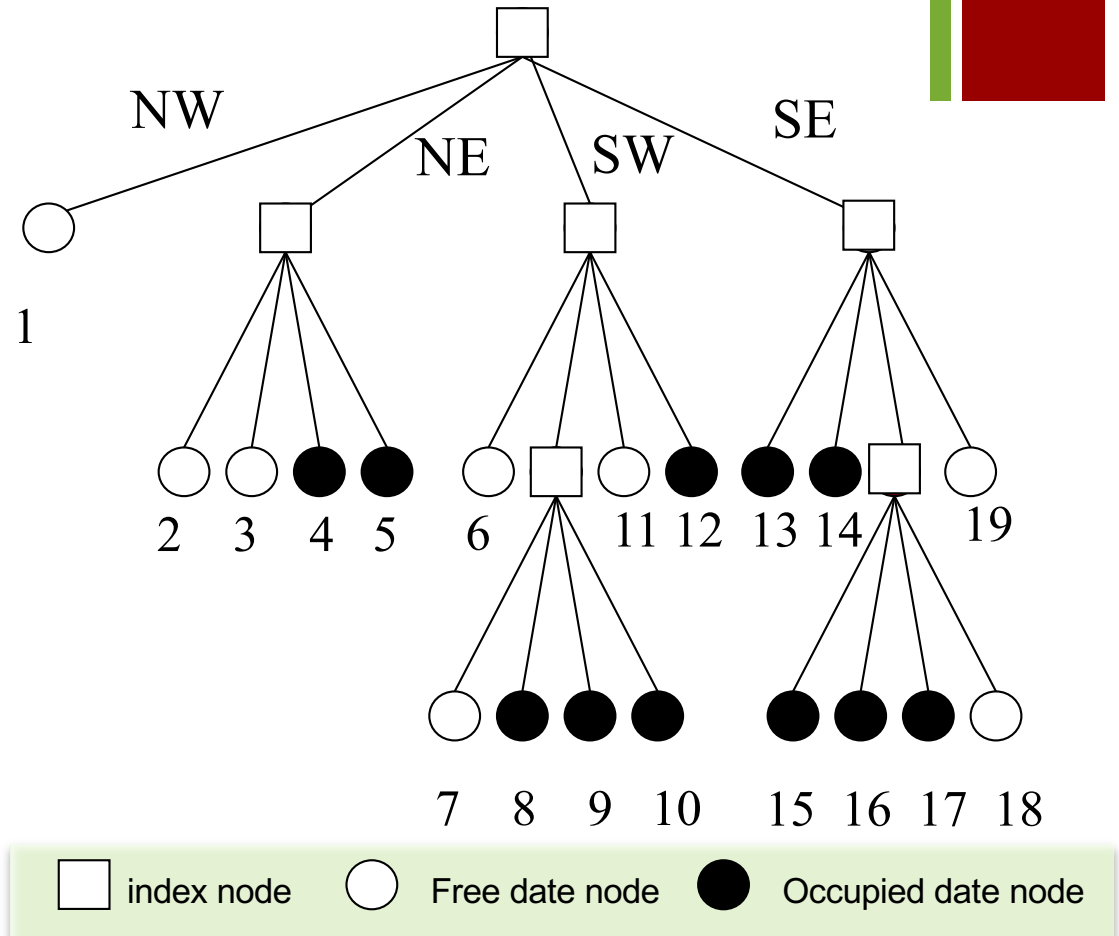


$n=29$ (25)

How many times do you need to fold a piece of paper to make it reach the moon?

Thickness of paper sheet = 0.1mm
Distance = 384,403km

+ Quadtree – Basic Idea

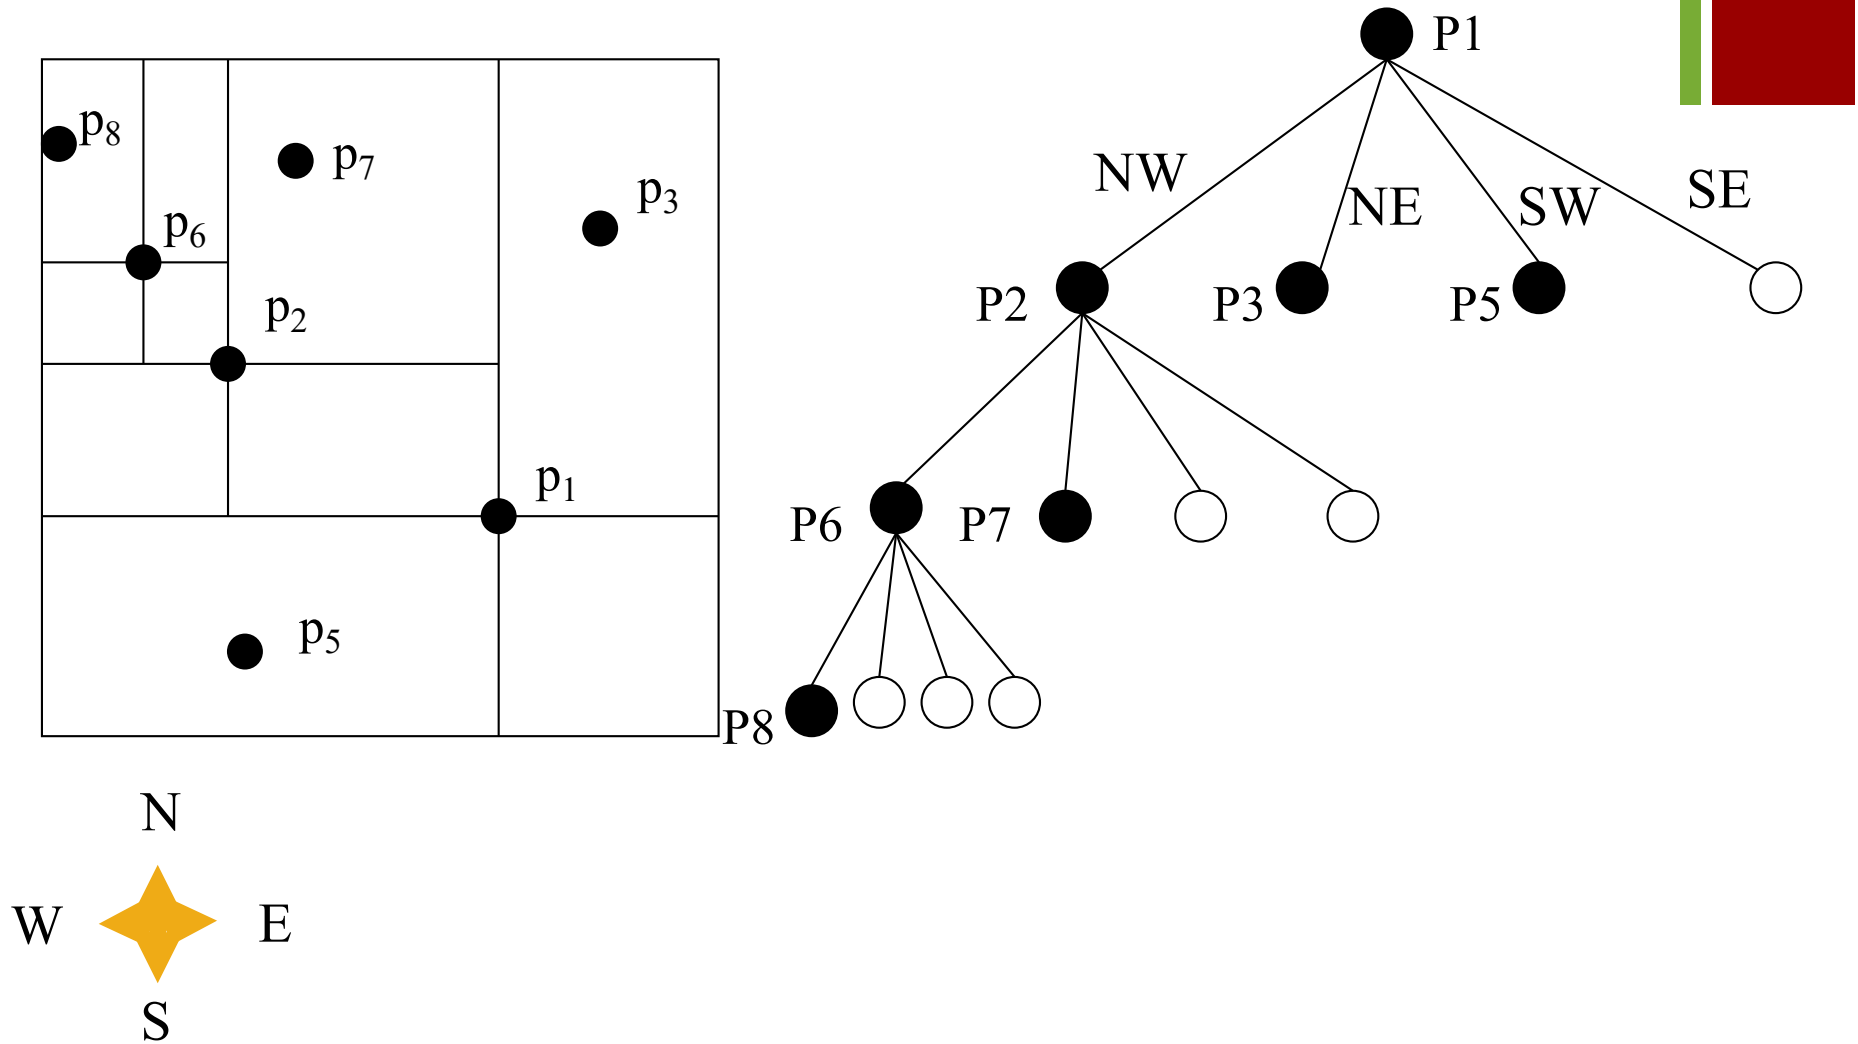
32



...originally proposed for compact image representation that supports image union/intersection operations efficiently

+ Point Quadtree Index



33



... the shape of the tree depends on the order to point insertion

+ Region Quadtree Index

34



... the shape of the tree does not depend on the order to point insertion

+ Some Questions

- Is the quadtree a balanced tree?
- When does deposition stop?
- What's the main difference between the point quadtree and the region quadtree?
- Will the order of point insertion affect the shape of a quadtree?
- What information a node needs to record?
- How to perform point and region queries using these two types of quadtrees?

+ kd-Tree

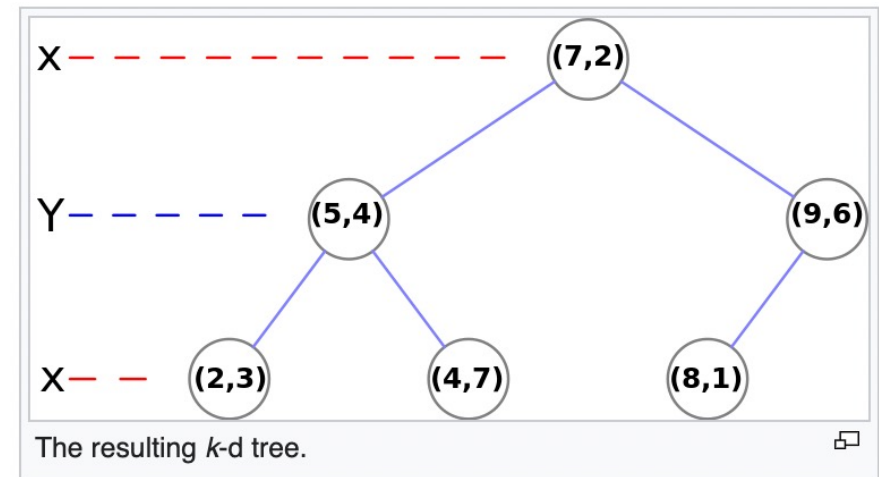
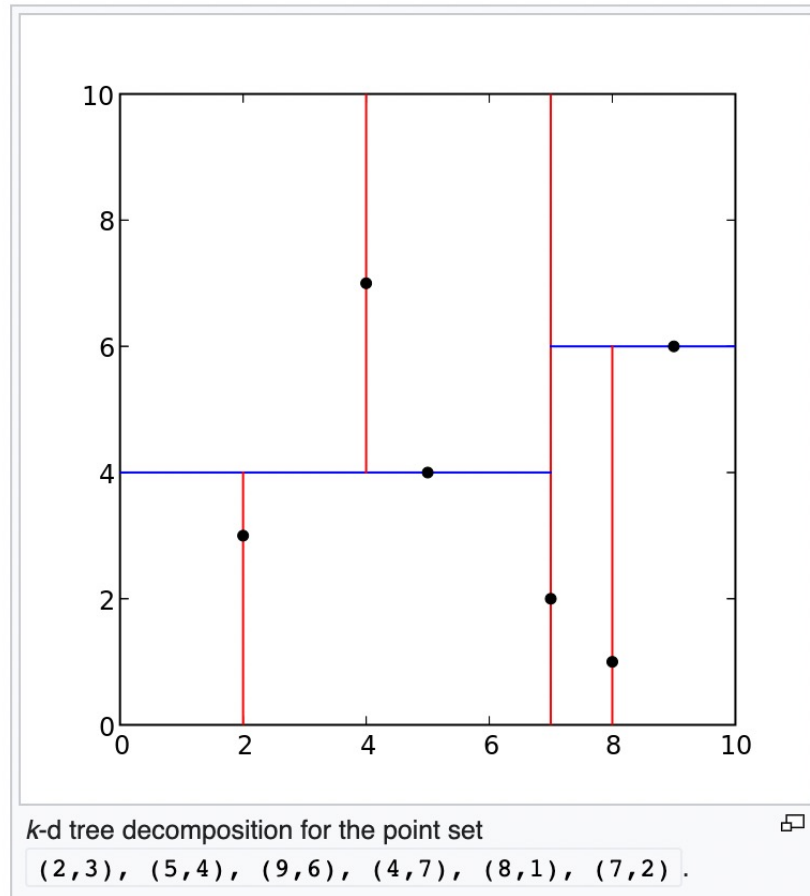
36

- Decomposition at data points (like the point quadtree)
- Motivation
 - For both point & region quadtrees: a node stores 2^k pointers and each time k comparisons are required to decide which subtree to go for a k -dimensional data
- Basic idea
 - Select **one** dimension, split according to this value and do the same **recursively** with the two new sub-partitions
 - **Fan-out** is a constant ($=2$) for arbitrary number of dimensions
 - Number of comparisons at each node is also a constant ($=1$)
- Problem
 - The resulting **binary tree** is high, and its *one-point-per-node structure* is not suitable for secondary storage

... very popular for in-memory data organization and commonly used in many application areas such robotics and computer vision

+ kd-Tree: Example

37



The tree structure is dependent on the order of insertion (not robust for sorted data)

Many variations: non-alternative, data at leaves only, balanced kd-trees, representing regions etc.

+ Some Questions


- How to use a kd-tree to support point query?
 - That is, to find if a given point exist in the database
 - Can you write the pseudo code for this?
- How to use a kd-tree to support range query?
 - That is, to find all points within a given rectangle?
 - Can you write the pseudo code for this?
- Later, we will discuss nearest neighbor queries, please come back to think about how you can support such a query using a kd-tree index
- What about insert or delete a point?
 - Typically used as an in-memory read-only data structure

+ Multidimensional Data

- There is no **total order** that preserves spatial proximity
- Solution:
 - Find heuristic solutions: use a total order mapping that can preserve the **spatial proximity** to some extent
- Idea:
 - If two objects are located close together in the original space (a k -dimensional), they should be close together in a target one-dimensional space (with high probability)
 - A balanced tree index structure for one-dimensional data is well known (B/B⁺ tree)

+ Space-Filling Curves



40



To find a space-filling curve in a high dimensional space to better preserve locality approximation



+ Row Order (or Column Order)

41



+ Z-Order (Peano Order)

42





- Easy and elegant way to encode cells
- SIRO-DBMS (SDM) and Oracle use this order

Giuseppe Peano (/pi'a:nov/) was an Italian mathematician (1858 – 1932)



+ Gray Order

43



Frank Gray was a physicist and researcher at Bell Labs (1887 – 1969)


+ Hilbert Order



David Hilbert was a German mathematician (1862 –1943).

+ Bit-Interleaving

45



- bit interleaving for calculating z-values

$x_0y_0x_1y_1\dots$

- an alternative way for calculating z-values
- works fine with varying resolutions

- example for cell A

$$(120)_4 = (24)_{10}$$

$$A_x=010, A_y=100$$

$$(011000)_2 = (24)_{10}$$

+ Some Properties of Z-Values

46

■ Variable length

- Approximate at different levels
- Appending '0's at the end to unify z-value length

■ Nesting Peano cells

- $a = a_1a_2a_3..a_n$
- $b = b_1b_2b_3...b_n$
- a is nested inside b if and only if
 - $\text{length}(a) \geq \text{length}(b)$, $\text{length}(a)$ is the number of non-zero digits in z-value a
 - let $k = \text{length}(a)$, $a_i = b_i$, $1 \leq i \leq k$

+ Using Z-Values and B-Tree (I)


47

■ Motivation

- Use standard B-tree to manage multidimensional data

■ Basic Idea

- a Peano cell corresponds to a bucket
- Peano cells are of varying sizes
- Z-values are managed by B-tree



+ Using Z-Values and B-Tree (II)

48


■ Search

- Point query: find the z-value for the unit Peano cell containing point p
- Range query: find the min and max z-values for rectangle R (or the z-values approximating R)

■ Insertion and deletion

■ Compatibility of z-value indices

- Origin and orientation
- Spatial extent
- Resolution



+ Data Access Methods

49

- One dimensional

- Hashing and B-Trees

- Point data

- Hashing: GRID and EXCELL
 - Hierarchical
 - Quadtree: point and region quadtrees
 - kd-Tree
 - Z-values and B-tree


- Polygon data

- Transformation: end point mapping and z-values
 - Overlapping: R-tree and R*-tree
 - Clipping: R+-tree

+ Indexing Objects with Spatial Extent

50

- Rectangles more difficult than points as they do not fall into a single cell of a bucket partition.
- Three strategies
 - Transformation
 - Overlapping regions
 - Clipping



+ Transformation: High Dimensional Points

51

■ Motivation

- Points are easy to manage


■ Basic Ideas

- A rectangle in 2-D space can be mapped to a point in 4-D space
- Using point access methods

■ Two methods


- Endpoint mapping, or midpoint mapping
 - $(x_{\text{low}}, y_{\text{low}}, x_{\text{high}}, y_{\text{high}}) \rightarrow (x_{\text{low}}, x_{\text{high}}, y_{\text{low}}, y_{\text{high}})$
 - $(x_{\text{center}}, y_{\text{center}}, x_{\text{ext}}, y_{\text{ext}}) \rightarrow (x_{\text{center}}, x_{\text{ext}}, y_{\text{center}}, y_{\text{ext}})$

+ Endpoint Mapping




+ Query Processing Using Endpoint Mapping

53



Given x-interval (l, u) :

- 1) *intersection query*: find all x-intervals **overlapping** with (l, u) .
- 2) *containment query*: find all x-intervals **inside** (l, u) .
- 3) *enclosure query*: find all x-intervals **enclosing** (l, u) .



+ Problems with Endpoint Mapping

54

- Points in the higher-D space are highly skewed
- Almost no relationship between the distances of two objects in the original space and the higher-D space
- A simple, intuitive query in the original space becomes complex and difficult to understand in the higher-D space
- Query processing in the higher-D space less efficient

...conceptually very simple, but...

+ Transformation: Using Z-Ordering

55

■ Motivation

- Still using point access methods, but without drawbacks of the previous approach

■ Basic idea

- Instead of mapping a polygon into a point, decompose a polygon into a set of Peano cells and map each Peano cell into a number (i.e., z-value)

+ Transformation: Using Z-Ordering


56

■ Granularity

- {11}, or {111, 112, 114}, or {111, 1121, 1123, 1124, 1141, 1142}

■ When decomposition stops

- Current cell either fully out or in the polygon
- Reached the “resolution”



...the entire space is 1.


+ Redundancy in Z-Ordering

57

■ Finer granularity

- ✓ Improves approximation accuracy
- ✓ Can reduce the number of “false hits”
- ✗ Too many index entries degrade query performance because of inflated index table
- ✗ May identify the same object multiple times in spatial query processing

■ The 4-Key method



+ Query Processing Using Z-Values

58

- Key techniques
 - Identification of nesting Peano cells by their z-values
 - Filter-and-refine
- Point query
- Window query
- Within buffer
- Spatial join query

+ Overlapping Regions

59

■ Motivation

- Single index entry for a polygon

■ Basic ideas

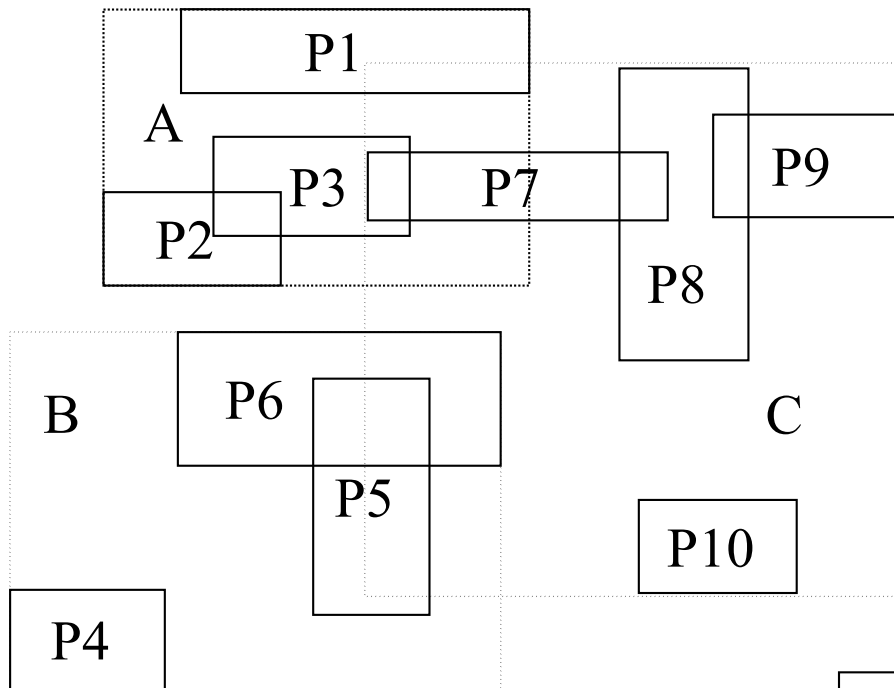
- One object (or its key) in one bucket only
- Cell boundary calculated according to polygons inside the cell
- Allow overlapping cells: *inevitable!*

■ Problems

- Multiple cells need to be examined to search an object
- Where to insert?

+ R-Tree and R*-Tree

60




A node must have more than m , less than M elements.

Many different strategies for:

- insertion \rightarrow split
- deletion \rightarrow reinsert

what info recorded in a node?



+ Query Processing Using R-Trees

61

- A node records the MBR of all objects in the subtree rooted from the node
- Point query
- Window query
- Within buffer
- Spatial join query

...home work

+ Clipping

62

■ Motivation

- Single search path for a point query

■ Basic ideas

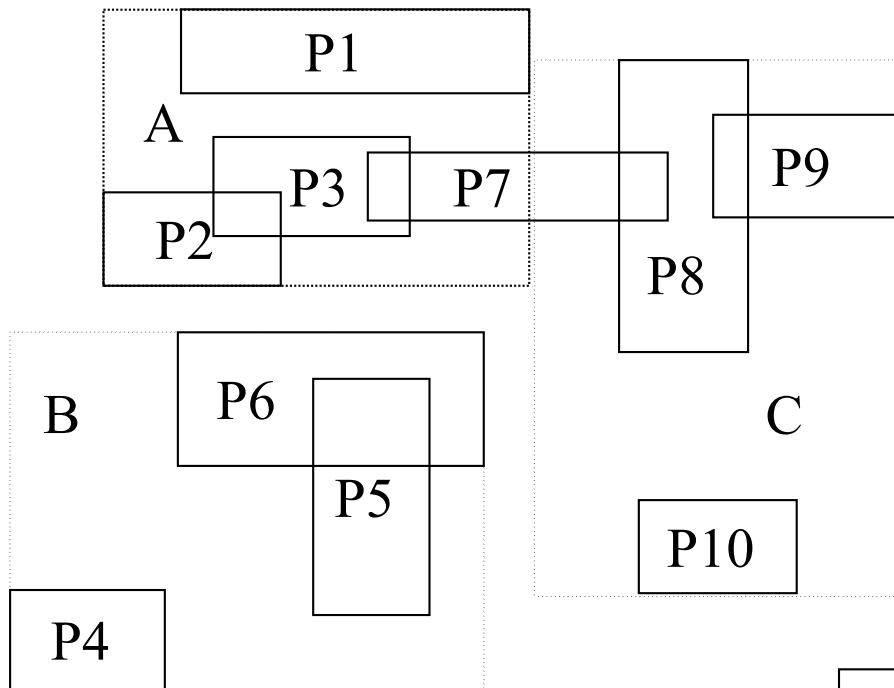
- Clipping polygon at cell boundaries
- Allowing one polygon in multiple cells

■ Problems

- Multiple index entries for an object (increased search time, overflow more likely)
- Enlargement deadlock
- Cascading splitting


+ R⁺-Tree

63



Insertion of an object can affect several nodes, and may cause 1) MBR enlargement; 2) deadlock; 3) node splitting (upwards and downwards, with object re-clipping).

Deletion may lead to merger, which is not always possible.



+ Query Processing Using R⁺-Trees

64

- Point query
- Window query
- Within buffer
- Spatial join query

...home work

+ Data Access Methods

65

■ One dimensional

- Hashing and B-Trees

■ Point data

- Hashing: GRID and EXCELL
- Hierarchical
 - Quadtree: point and region quadtrees
 - kd-Tree
 - Z-values and B-tree

■ Polygon data

- Transformation: end point mapping and z-values
- Overlapping: R-tree and R*-tree
- Clipping: R+-tree

+ Indexing High Dimensional Data

- GIS applications in 2- or 3-D only
- Multimedia DB can have data with several hundred dimensions.
- While point/polygon access methods can be generalized for higher-D applications, they may be not efficient
- High-D indexing is a hard problem