

CSIT6000P Spatial and Multimedia Databases
2022 Spring

RDBMS EXercises

Prof Xiaofang Zhou

+ Assessment

- Two assessment components
 - Two assignments (individual): 30% + 30% → out soon
 - Marking criteria can be found in the assignment specs
 - Final report and presentation (group): 40% → now out
 - Marking criteria can be found in the project spec
- No exams
- Your final score for this course is the sum of the above components

+ DBMS

- What are the main components of a database system?
- What are the main functions of a RDBMS?
- Name some RDBMS products

+ Dimensions

- What data types can a RDBMS support?
- Compare and contrast multi-attributes data and multi-dimensional data
- Identify applications with one, several and many dimensions
- Why a RDBMS cannot support spatial data (such as points, lines, and polygons)?
- Discuss linear scan based search and tree-based search for the above datasets

+ DB Design

- Database normalization

- 1st, 2nd and 3rd normal forms: definitions and issues
- Join and denormalization

- Integrity constraints and consistency

- Data type constraints, system constraints, app constraints

- Three schemas and data independence

+ Query Execution

- SQL and query execution planning
- Database transactions

+ Data Types

- What data types can an RDBMS support?
- What does a data type mean?

+ SQL

- SQL = DDL + DML + DCL
- DDL: CREATE, ALTER and DROP tables
- DML: SELECT, INSERT, DELETE and UPDATE
- DCL: GRANT, REVOKE

+ SELECT

SELECT *a list of expressions*

FROM *a list of tables*

WHERE *a list of conditions (AND, OR)*

GROUP BY *a list of expressions to be selected*

HAVING *conditions defined on groups*

ORDER BY *a list of expressions to be selected*

+ SELECT Query Example

10

To find the departments and their managers who are female of 30 years old or younger.

```
SELECT      e1.firstName, UPPER(e1.lastName), d.name
FROM        Department d, Employee e1, Employee e2
WHERE       d.id = e.dNum AND e1.ssn = e2.mgrSSN AND
              e1.sex = 'female' AND
              DATEDIFF(YEAR, e1.dob, GETDATE()) < 30;
```

```
Employee[fName, lName, ssn, dob, add, sex, salary, mgrSSN, dNum]
```

```
Department[name, id, mgrSSN, mgrStartDate]
```

+ SELECT With Sub-query

11

To find the departments **with salary budget over 10M** and their managers who are female of 30 years old or younger.

```
SELECT e1.firstName, UPPER(e1.lastName), d.name
FROM   Department d, Employee e1, Employee e2
WHERE  d.id = e1.dNum AND e1.ssn = e2.mgrSSN AND
       e1.sex = 'female' AND
       DATEDIFF(YEAR, e1.dob, GETDATE()) < 30 AND
       d.id = (SELECT      id
                FROM        Department, Employee
                WHERE        id = dNum
                GROUP BY    id
                HAVING      SUM(salary) >= 100000000);
```

```
Employee[fName, lName, ssn, dob, add, sex, salary, mgrSSN, dNum]
```

```
Department[name, id, mgrSSN, mgrStartDate]
```

+ QEP

- SQL queries are **declarative**
- An SQL query will be translated into an execution plan by the QEP
 - One SQL query can have many **equivalent** execution plans
 - But they can have **different** execution costs
- The QEP selects the “**best**” execution plan
 - Optimization can be based on **cost models**, or **heuristics**
 - QEP requires data **statistics** and **index** information
 - Common heuristics: **push-down selections**
 - Note the pros and cons of **pushing-down projections**

+ QEP Discussion

13

To find the departments and their managers who are female of 30 years old or younger.

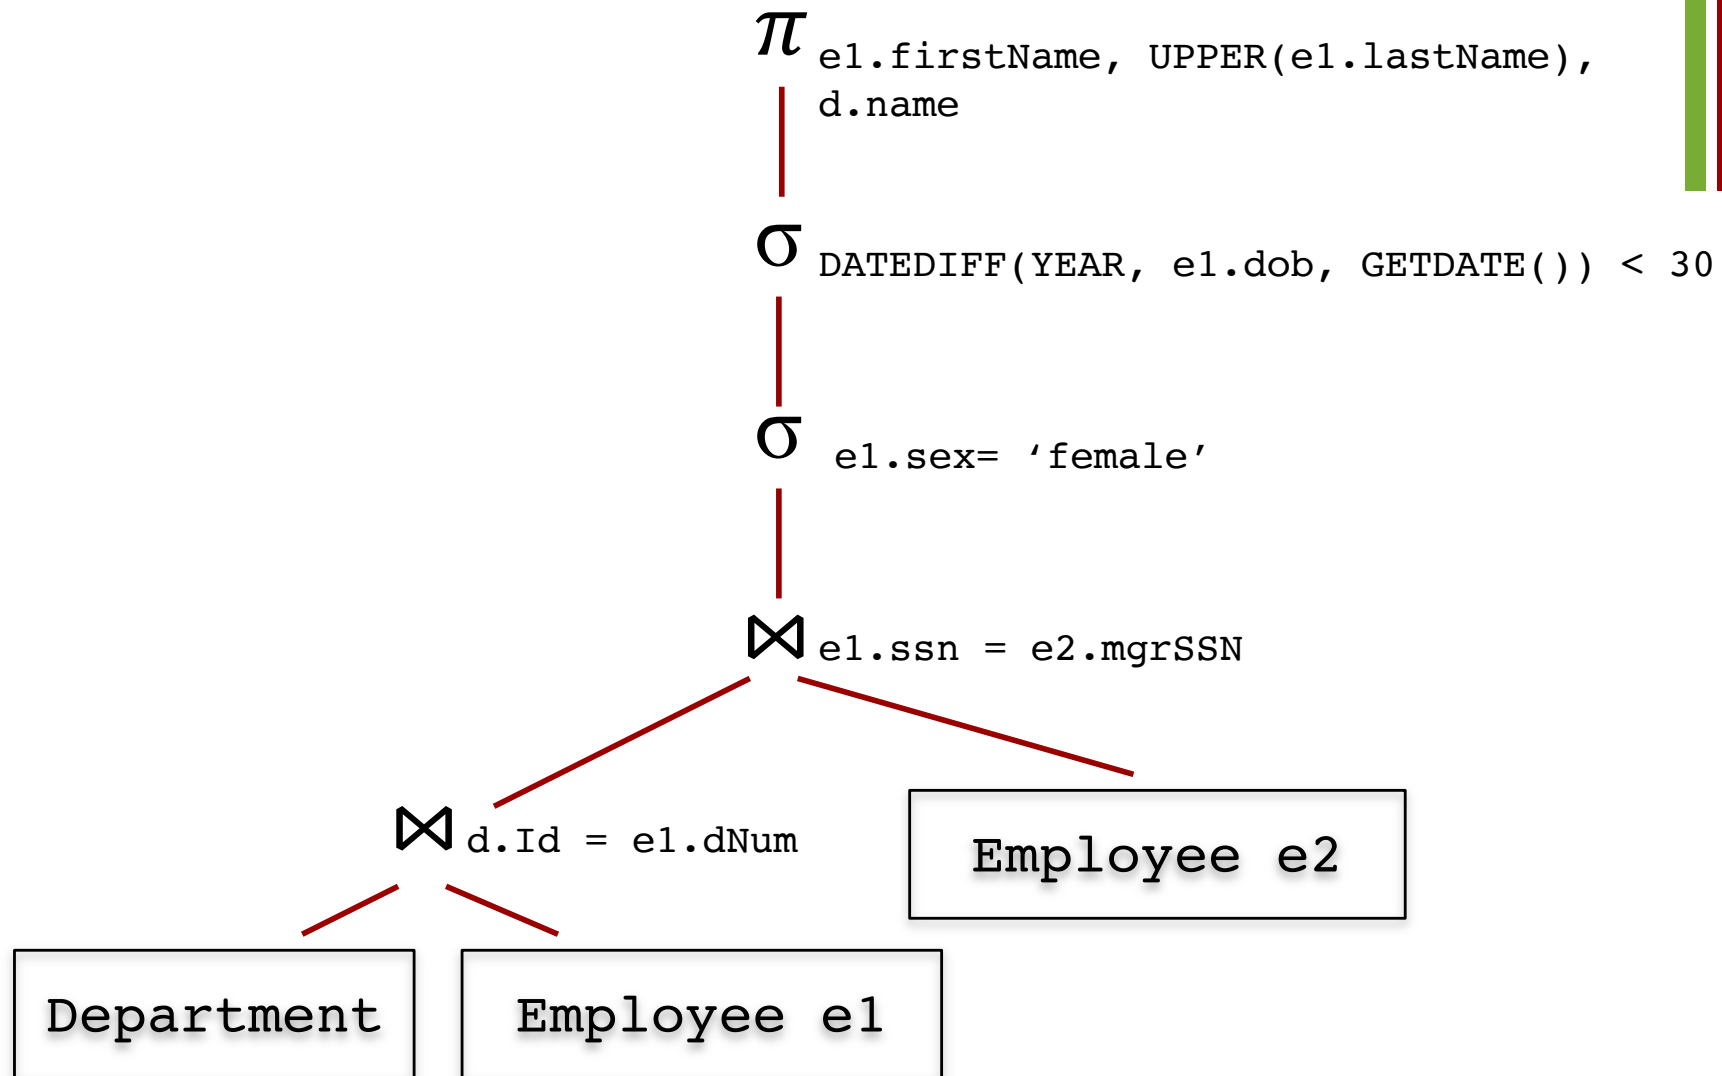
```
SELECT      e1.firstName, UPPER(e1.lastName), d.name
FROM        Department d, Employee e1, Employee e2
WHERE        d.id = e1.dNum AND e1.ssn = e2.mgrSSN AND
              e1.sex = 'female' AND
              DATEDIFF(YEAR, e1.dob, GETDATE()) < 30;
```

```
Employee[fName, lName, ssn, dob, add, sex, salary, mgrSSN, dNum]
```

```
Department[name, id, mgrSSN, mgrStartDate]
```

+ An Execution Tree: Straightforward

14

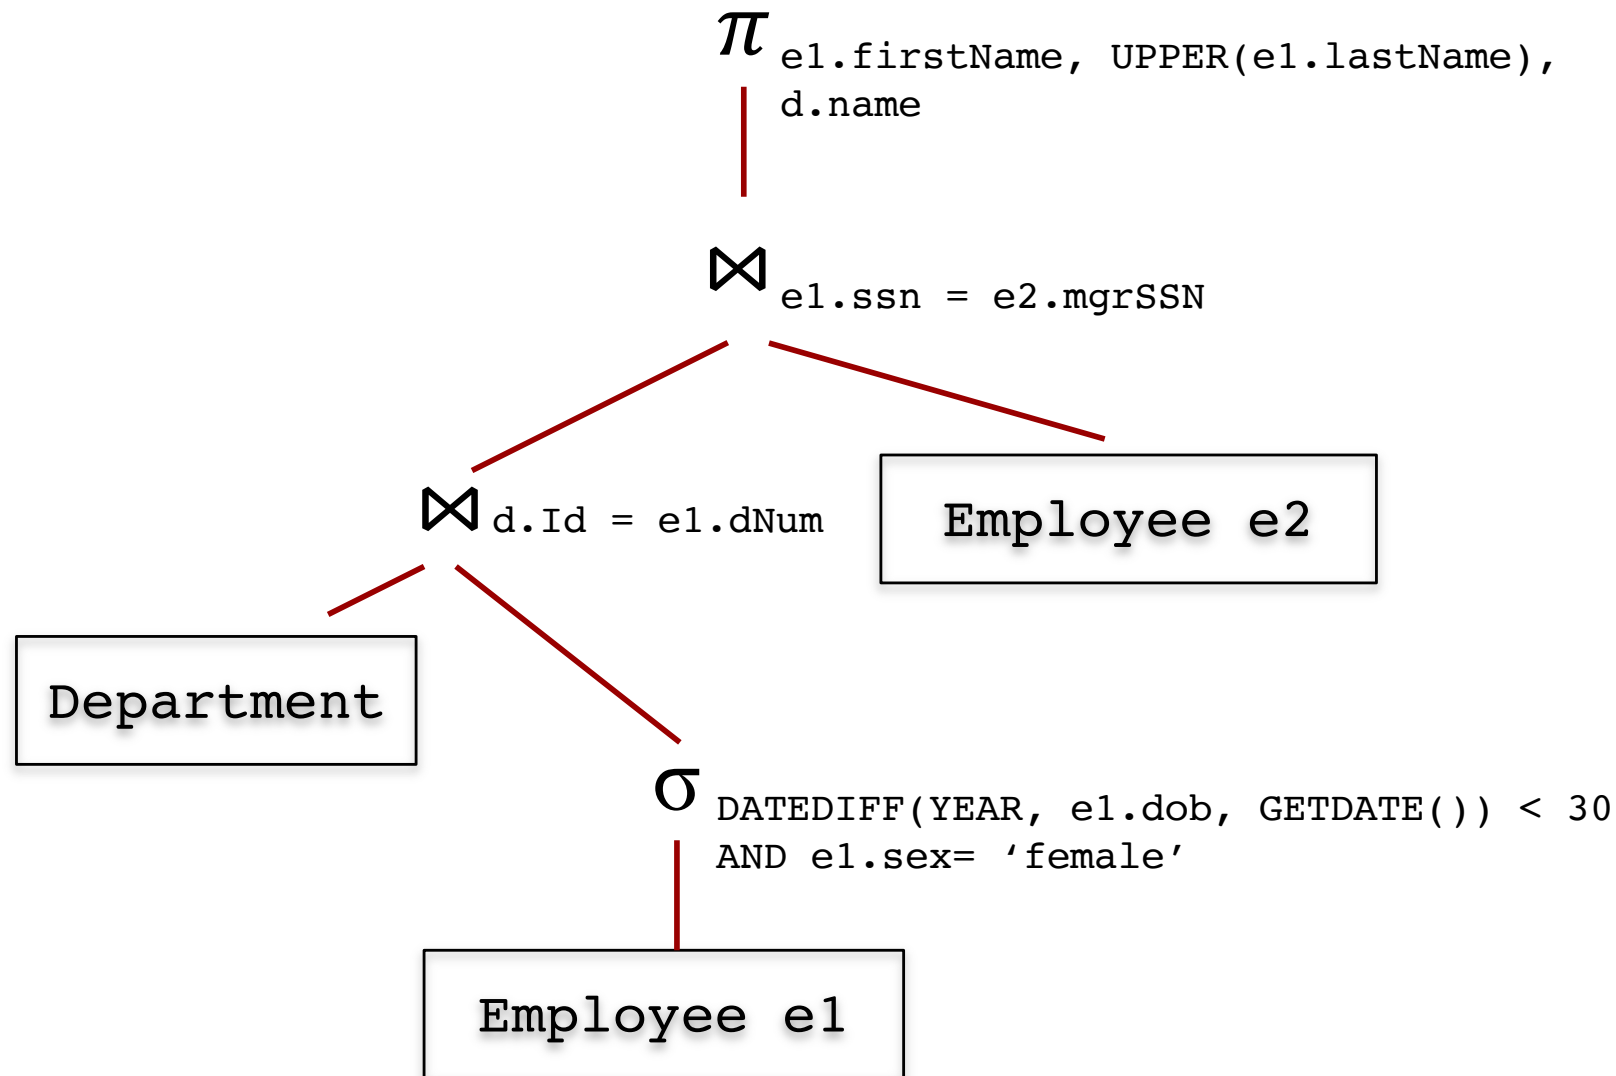


Employee[fName, lName, ssn, dob, add, sex, salary, mgrSSN, dNum]

Department[name, id, mgrSSN, mgrStartDate]

+ An Execution Tree: Better

15

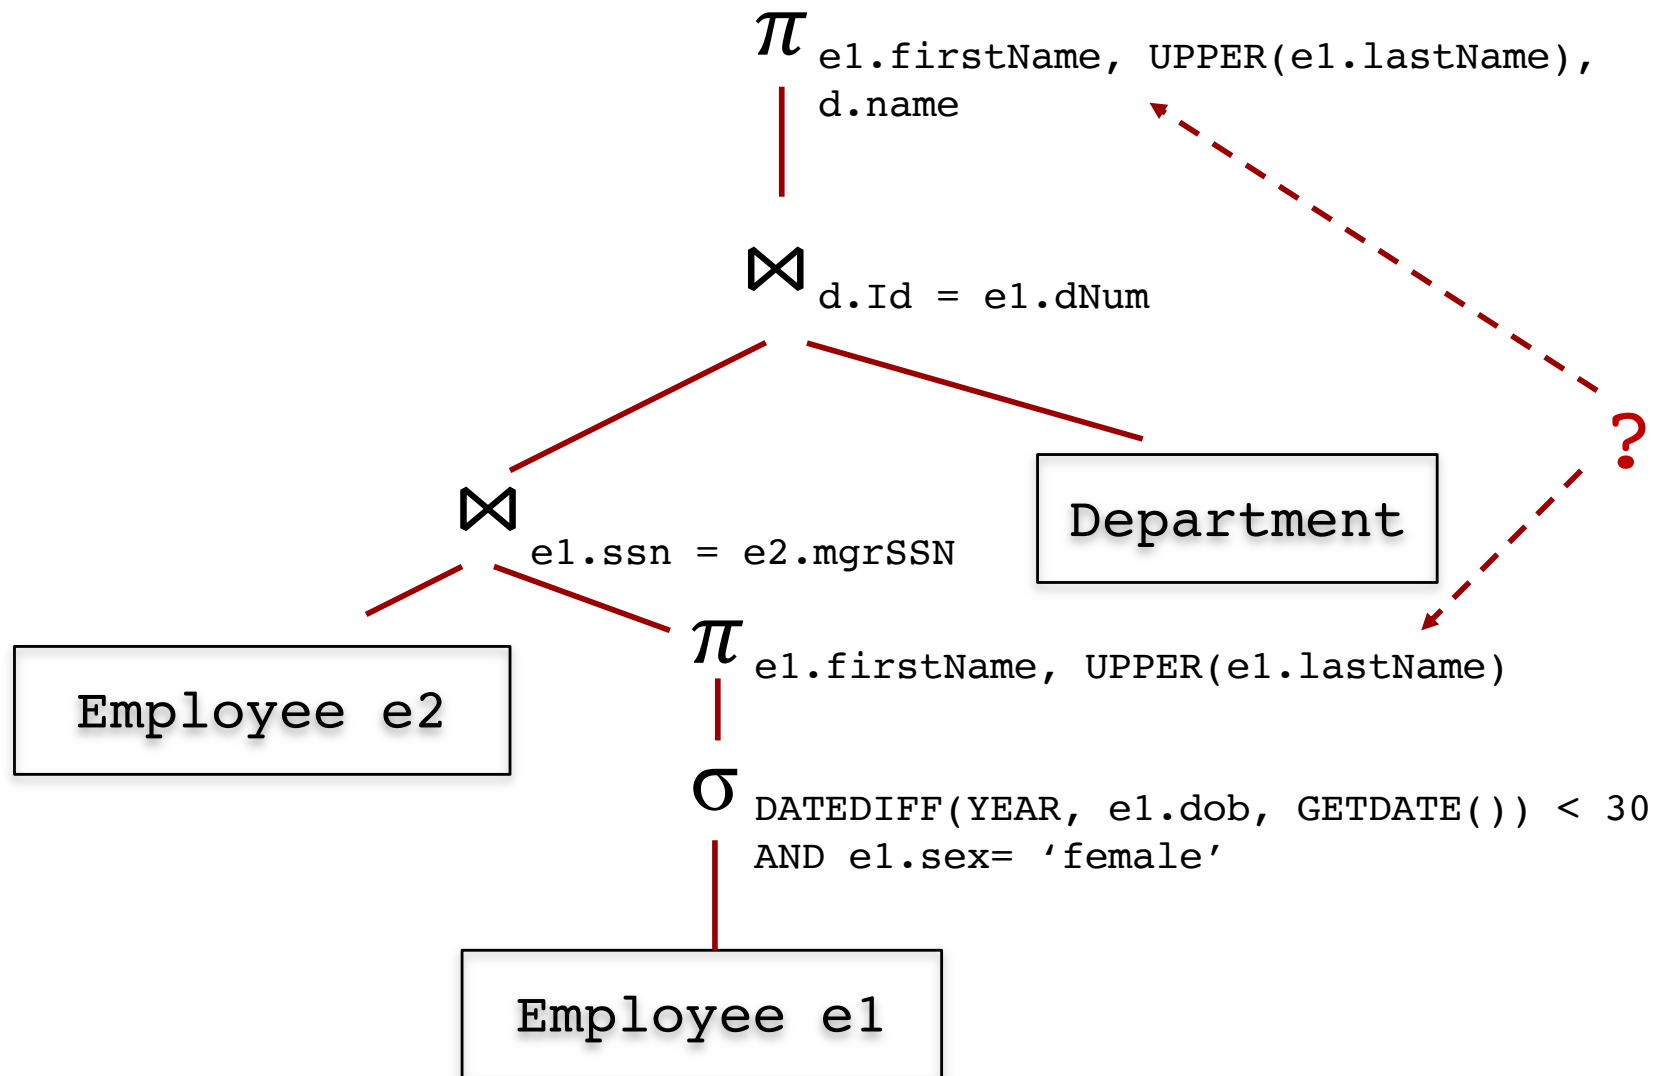


Employee[fName, lName, ssn, dob, add, sex, salary, mgrSSN, dNum]

Department[name, id, mgrSSN, mgrStartDate]

+ An Execution Tree: Even Better?

16



Employee[fName, lName, ssn, dob, add, sex, salary, mgrSSN, dNum]

Department[name, id, mgrSSN, mgrStartDate]

+ Indexing

- B-tree and B⁺-tree

- Bitmap index

- Hash index

+ RDBMS Summary

18

- RDBMS functions

- Define and construct tables, data manipulation, data maintenance

- Database design

- Conceptual, logical and physical design, indexing, data access control (DAC)

- Integrity constraints

- SQL = DDL + DML + DCL

- Query optimization

- ACID properties

- Atomicity, consistency, isolation, durability